



UMD DATA605 - Big Data Systems

8.3: Apache Hadoop

- **Instructor:** Dr. GP Saggese - gsaggese@umd.edu
- **References**
 - Ghemawat et al.: *The Google File System*, 2003
 - Dean et al.: *MapReduce: Simplified Data Processing on Large Clusters*, 2004

Hadoop Ecosystem (aka Hadoop Zoo)

- **Hadoop MapReduce**
 - A framework for processing large data sets in parallel across a Hadoop cluster
 - **HDFS**
 - Distributed file system
 - **Pig**
 - High-level data-flow framework for parallel computation
 - **HBase**
 - Scalable, distributed database
 - Structured data storage for large tables (like Google BigTable)
 - **Cassandra**
 - Scalable multi-master database with no single points of failure
 - **Hive**
 - Data warehouse infrastructure
 - Provide data summarization and ad-hoc querying
 - **ZooKeeper**
 - High-performance coordination service for distributed applications
- YARN, Kafka, Storm, Spark, Solr, ...**



Hadoop Distributed File System (HDFS)

- HDFS is a **distributed file system**
 - Designed to store large data sets reliably
 - Part of the Apache Hadoop ecosystem
 - Inspired by the Google File System (GFS)



1. Optimized for **high-throughput access** to large files
 - Suitable for batch processing
 - Not low-latency access

2. Designed for **fault tolerance and scalability**
 - Ensures fault tolerance through replication
 - Blocks are stored on different nodes and racks
 - Provides data availability even if some nodes fail
 - Follows a primary-secondary architecture
 - Replication strategy improves read performance

HDFS Architecture

- **NameNode**


- Store file/dir hierarchy
- Store file metadata
 - E.g., block location, size, permissions

- **DataNodes**

- Store actual data blocks
- Split file into 16-256MB blocks
- Replicate chunks (2x or 3x) across multiple *DataNodes*
- Keep replicas in different racks

- **Client**

- API (e.g., Python, Java) to library
- Mount HDFS on local filesystem



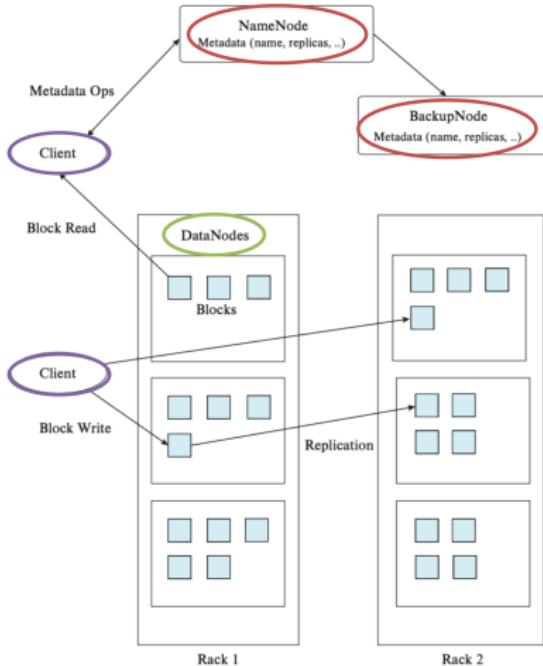
HDFS: Read / Write Protocols

- **Read**

- Contact *NameNode* for *DataNode* and block pointer
- Choose nearest *DataNode* for each block
- Connect to *DataNode* for data access
- Read blocks in parallel to improve performance
- Client reassembles data in correct order

- **Write**

- *NameNode* creates blocks
- Assign blocks to multiple *DataNodes*
- Client sends data to *DataNodes*
- *DataNodes* store data
- Pipeline blocks to other replicas
- Write successful after all replicas acknowledge



Fault Tolerance and Recovery

- *NameNode* monitors *DataNode* heartbeat signals
 - On failure, blocks are re-replicated to maintain replication factor
- *NameNode* itself is a single point of failure
 - Solved with HDFS High Availability
- Data integrity ensured using checksums

HDFS vs Traditional File Systems

- Best for **storing and processing large-scale files**
 - E.g., logs, media, sensor data
 - Commonly used in data lakes and ETL pipelines
 - Supports very large files and directories
 - Performance degrades with many small files
- Optimized for **write-once, read-many** access pattern
- Lacks low-latency access, but provides **high throughput**
 - Good for analytics (OLAP)
 - Not suitable for transactional systems (OLTP)
 - E.g., bank

MapReduce: Hadoop

- **Hadoop**: open-source MapReduce implementation



- **Functionalities**

- Partition input data (HDFS)
- Input adapters
 - E.g., HBase, MongoDB, Cassandra, Amazon Dynamo
- Schedule program execution across machines
- Handle machine failures
- Manage inter-machine communication
- Perform *GroupByKey* step
- Output adapters
 - E.g., Avro, ORC, Parquet
- Schedule multiple *MapReduce* jobs