


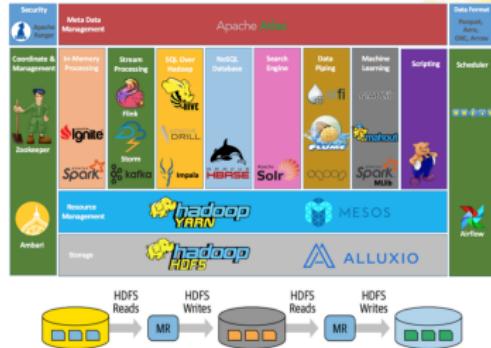
9.1: Apache Spark: Principles

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu
- **References:**
 - Concepts in the slides
 - Academic paper
 - “Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing”, 2012
 - Mastery
 - “Learning Spark: Lightning-Fast Data Analytics” (2nd Edition)
 - Not my favorite, but free



Hadoop MapReduce: Shortcomings

- **Hadoop is hard to administer**
 - Many layers (HDFS, Yarn, Hadoop, ...)
 - Extensive configuration
- **Hadoop is hard to use**
 - Verbose API
 - Limited language support (e.g., Java is native)
 - MapReduce jobs read / write data on disk
- **Large but fragmented ecosystem**
 - No native support for:
 - Machine learning
 - SQL
 - Streaming
 - Interactive computing
 - New systems developed on Hadoop for new workloads
 - E.g., Apache Hive, Storm, Impala, Giraph, Drill



(Apache) Spark



- Open-source
 - DataBrick monetizes it (\$100B startup in 2025)
- General processing engine
 - Large set of operations beyond Map() and Reduce()
 - Combine operations in any order
 - Computation organized as a DAG, decomposed into parallel tasks
 - Scheduler/optimizer for parallel workers
- Supports several languages
 - Java, Scala (preferred), Python supported through bindings
- Data abstraction
 - Resilient Distributed Dataset (RDD)
 - DataFrames, Datasets built on RDDs
- Fault tolerance through RDD lineage
- In-memory computation

Keep intermediate results in memory, if possible


Berkeley: From Research to Companies

- Pathway from lab innovation to startups
 - Students and researchers creating companies from lab systems
 - Focus on data-intensive systems and machine learning
 - Open-source ecosystems enabling broad adoption
- **AMPLab**
 - Collaborative projects creating systems like Spark
 - Industry engagement guiding real-world impact
- **RISELab**
 - Shift to systems supporting AI, security, and automation
 - Platforms like Ray and ML-focused infrastructure



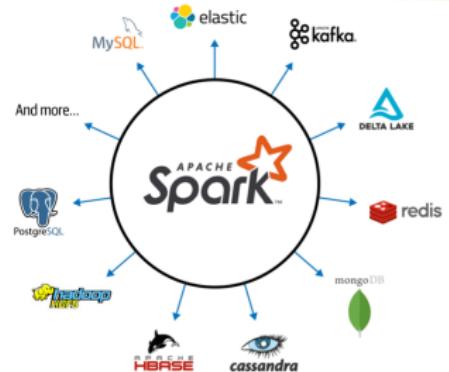
Berkeley AMPLab Data Analytics Stack

- So many tools that they have their own Big Data stack!
<https://amplab.cs.berkeley.edu/software/>

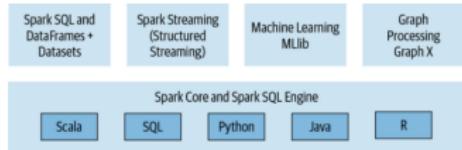


Apache Spark: Introduction

- **Unified stack**
 - Different computation models in a single framework
 - **Spark SQL**
 - ANSI SQL compliant
 - Work with structured relational data
 - **Spark MLLib**
 - Build ML pipelines
 - Support popular ML algorithms
 - Built on Spark DataFrame
 - **Spark Streaming**
 - Handle continually growing tables
 - Treat tables as static
 - **GraphX**
 - Manipulate graphs
 - Perform graph-parallel computation
- **Extensibility**
 - Read from many sources
 - Write to many backends



One computation engine



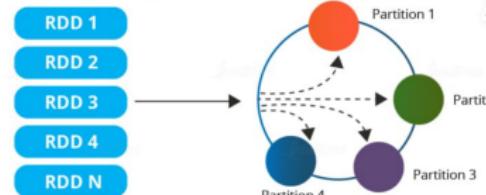
General purpose applications



Resilient Distributed Dataset (RDD)

- **Resilient Distributed Dataset (RDD)**

- Collection of data elements
- Partitioned across nodes
- Operated on in parallel
- Fault-tolerant
- In-memory / serializable



- **Applications**

- Best for applying the same operation to all dataset elements (vectorized)
- Less suitable for asynchronous fine-grained updates to shared state
 - E.g., updating one value in a dataframe

- **Ways to create RDDs**

- Reference data in external storage
 - E.g., file-system, HDFS, HBase
- Parallelize an existing collection in your driver program
- Transform RDDs into other RDDs


Transformations vs Actions

- **Transformations**

- Lazy evaluation
- Compute only when an Action requires it
- Build a graph of transformations

- **Actions**

- Aka “materialize”
- Force calculations on RDDs and return values



Spark Example: Estimate Pi

- Goal

- Estimate π using random sampling in the unit square
- Fraction of points inside the unit circle approximates $\pi/4$

- sample generates one random point
 - Test membership in the unit circle
 - Returns 1 for inside, 0 for outside

- parallelize distributes the sampling task
 - Each element in the RDD triggers one call to sample
 - “Embarrassingly parallel” computation
- map applies sampling across partitions
 - Each worker independently counts hits inside the circle
- reduce aggregates partial sums
 - Summing 0 and 1 values yields total count of hits

```
# Estimate pi (compute-intensive task).
# Pick random points in the unit square [(0,0)-(1,1)].
# See how many fall in the unit circle center=(0, 0), radius=1.
# The fraction should be pi / 4.

import random
random.seed(314)

def sample(p):
    x, y = random.random(), random.random()
    in_unit_circle = 1 if x*x + y*y < 1 else 0
    return in_unit_circle

# "parallelize" method creates an RDD.
NUM_SAMPLES = int(1e6)
count = sc.parallelize(range(0, NUM_SAMPLES)) \
    .map(sample) \
    .reduce(lambda a, b: a + b)
approx_pi = 4.0 * count / NUM_SAMPLES
print("pi is roughly %f" % approx_pi)
executed in 386ms, finished 04:27:53 2022-11-23
pi is roughly 3.141400
```

Spark: Architecture

- **Architecture**

- Who does what
- I.e., responsibilities of each component

- **Spark Application**

- Code describing computation
- E.g., Python code calling Spark

- **Spark Driver**

- Transform operations into DAG computations
- Distribute task execution across *Executors*
- Communicate with *Cluster Manager* for resources

- **Spark Session**

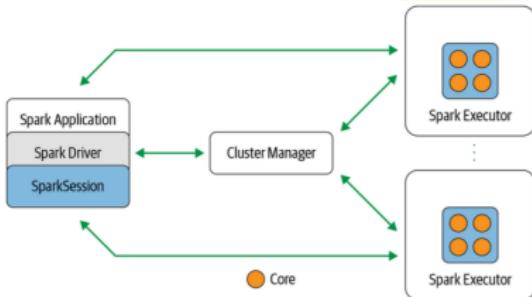
- Interface to Spark system

- **Cluster Manager**

- Manage and allocate resources
- Support Hadoop, YARN, Mesos, Kubernetes


- **Spark Executor**

- Run worker node to execute tasks
- Typically one executor per node
- Relies on JVM



Spark: Computation Model


- **Architecture**
 - Who does what
- **Computation model**
 - How are things done



- **Spark Driver**
 - Converts *Application* into *Jobs*
 - Describes computation with *Transformations* and triggers with *Actions*
- **Spark Job**
 - Parallel computation in response to a *Action*
 - Each *Job* is a DAG with dependent *Stages*
- **Spark Stage**
 - Smaller operation within a *Job*
 - Stages run serially or in parallel
- **Spark Task**
 - Each *Stage* has multiple *Tasks*
 - Single unit of work sent to a *Executor*
 - Each *Task* maps to a single core and works on a single data partition


Distributed Data and Partitions

- **Data is distributed** as partitions across physical nodes
 - Store each partition in memory
 - Enable efficient parallelism
- **Spark Executors** process data “close” to them
 - Minimize network bandwidth
 - Ensure data locality
 - Similar to Hadoop



Parallelized Collections

- Parallelized collections created by calling *SparkContext parallelize()* on an existing collection



- Data spread across nodes
- Number of *partitions* to cut dataset into
 - Spark runs one *Task* per partition
 - Aim for 2-4 partitions per CPU
 - Spark sets partitions automatically based on your cluster
 - Set manually by passing as a second parameter to `parallelize()`

Deployment Modes

- Spark can run on several different configurations
 - Components (e.g., Driver, Cluster Manager, and Executors) split on different nodes

Deployment Mode	Where Components Run	Notes
Local	Run in a single JVM on one machine	Run Spark on a laptop
Standalone	Run in separate JVMs on different machines	Spark's built-in cluster manager
YARN / Kubernetes	Run in different pods/containers	Production clusters