


9.1: Apache Spark: Primitives


- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu
- **References:**
 - Key concepts discussed in the slides
 - Academic paper
 - “Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing”, 2012
 - Mastery
 - “Learning Spark: Lightning-Fast Data Analytics” (2nd Edition)
 - Not my favorite, but free



Transformations vs Actions

- **Transformations**

- Transform a Spark RDD into a new RDD without modifying input data
- Immutability like functional programming
- E.g., `select()`, `filter()`, `join()`, `orderBy()`



- Transformations are **evaluated lazily**

- Spark optimizes the computation by analyzing the workload
 - E.g., joining, pipeline operations, breaking into stages
- Record results as "lineage"
 - The sequence of stages is rearranged and optimized without changing the results

- **Actions**

- Trigger computation evaluation
- E.g., `show()`, `take()`, `count()`, `collect()`, `save()`

Spark Example: MapReduce in 1 or 4 Line

```
|more data.txt
```

executed in 1.77s, finished 04:37:35 2022-11-23

One a penny, two a penny, hot cross buns

```
lines = sc.textFile("data.txt").flatMap(lambda line: line.split(" "))
pairs = lines.map(lambda s: (s, 1))
counts = pairs.reduceByKey(lambda a, b: a + b)
result = counts.collect()
print(result)
```

executed in 428ms, finished 04:36:24 2022-11-23

```
[('One', 1), ('two', 1), ('hot', 1), ('cross', 1), ('a', 2), ('penny,', 2), ('buns', 1)]
```

MapReduce in 4 Spark lines

```
result = sc.textFile("data.txt").flatMap(lambda line: line.split(" ")).map(
    lambda s: (s, 1)).reduceByKey(lambda a, b: a + b).collect()
print(result)
```

executed in 591ms, finished 05:06:00 2022-11-23

```
[('One', 1), ('two', 1), ('hot', 1), ('cross', 1), ('a', 2), ('penny,', 2), ('buns', 1)]
```

MapReduce in 1 (show-off) line

Same Code in Java Hadoop

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }
    }
}
```

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
                      Context context
                      throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
```



Spark Example: Logistic Regression in MapReduce

- Logistic Regression

```
# Load points
```

```
points = spark.textFile(...).map(parsePoint).c
```

```
# Initial separating plane
```

```
w = numpy.random.ranf(size=D)
```

```
# Until convergence
```

```
for i in range(ITERATIONS):
```

```
    # Parallel loop over the samples i=1...m
```

```
    gradient = points.map(
```

```
        lambda p:
```

```
            (1 / (1 + exp(-p.y*(w.dot(p.x)))) -  
             p.y * p.x
```

```
).reduce(lambda a, b: a + b)
```

```
w -= alpha * gradient
```

```
print("Final separating plane: %s" % w)
```

Repeat until convergence {

$$\theta_j \leftarrow \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

```
}
```

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2$$

Repeat {

$$\theta_j := \theta_j - \frac{\alpha}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

```
}
```

Spark Transformations: 1 / 3

- `map(func)`
 - Return a new RDD by applying `func()` to each element
- `flatMap(func)`
 - Map each input item to 0 or more output items
 - `func()` returns a sequence
- `filter(func)`
 - Return a new RDD selecting elements where `func()` returns true
- `union(otherDataset)`
 - Return a new RDD with the union of elements in the source dataset and the argument
- `intersection(otherDataset)`
 - Return a new RDD with the intersection of elements in the source dataset and the argument

From [here](#)

Spark Transformations: 2 / 3


- `join(otherDataset, [numTasks])`
 - On RDDs (K, V) and (K, W), return dataset of (K, (V, W)) pairs for each key
 - Support outer joins: `leftOuterJoin`, `rightOuterJoin`, `fullOuterJoin`
- `groupByKey([numPartitions])`
 - On RDD of (K, V) pairs, returns (K, Iterable<V>) pairs
 - For aggregation (e.g., sum, average), use `reduceByKey` for better performance
 - Process data in place instead of iterators
 - Output parallelism depends on parent RDD partitions
 - Use `numPartitions` to set tasks
- `sortByKey([ascending], [numPartitions])`
 - Returns (K, V) pairs sorted by keys in ascending or descending order

Spark Actions

- `reduce(func)`
 - Aggregate dataset elements using `func()`
 - `func()` takes two arguments and returns one
 - `func()` must be commutative and associative for parallel computation
- `collect()`
 - Return dataset elements as an array
 - Useful after operations producing a small data subset (e.g., `filter()`)
- `count()`
 - Return the number of elements in the dataset
- `take(n)`
 - Return an array with the first `n` dataset elements
 - `.collect()[:n]` differs from `.take(n)`
- From [here](#)


Spark: Fault-tolerance

- Spark leverages *immutability* and *lineage* for fault tolerance
- In case of **failure**
 - Reproduce RDD by replaying the lineage
 - Checkpoints aren't needed
 - Keep data in memory for increased performance
- Fault-tolerance is free!




Spark: RDD Persistence

- **Users explicitly cache an RDD**
 - I.e., `persist()`, `unpersist()`
 - Cache if RDD is expensive to compute
 - E.g., filtering large data
 - When you persist an RDD, each node:
 - Stores partitions of the RDD (in memory or on disk)
 - Reuses cached partitions on derived datasets
- **Cache**
 - Enhances speed of future actions (often by >10x)
 - Managed by Spark using LRU policy + garbage collector
- **Users can choose the storage level**
 - `MEMORY_ONLY` (default)
 - `DISK_ONLY` (e.g., Python Pickle)
 - `MEMORY_AND_DISK`
 - If an RDD doesn't fit in memory, store it on disk
 - Caching on disk can be more costly than not caching
 - Caching everything is often a bad idea



Spark: RDD Persistence and Fault-tolerance


- Spark unifies persistence and fault-tolerance through RDD lineage
- **Caching and Persistence**
 - Store RDDs in memory or disk to avoid repeated computation
 - Useful for iterative algorithms and interactive queries
 - E.g., `cache()` (memory only) or `persist()` (custom storage level)



- **Fault-Tolerance Mechanism**
 - RDDs are immutable and record lineage of transformations
 - If a partition is lost, Spark reconstructs it using lineage
 - No need to checkpoint unless lineage is too long
- **Persistence is Fault-Tolerant**
 - Cached RDDs can be recomputed from lineage if data is lost
 - Ensures recovery without manual intervention

Spark Shuffle

- Certain Spark operations trigger a data shuffle
- E.g., **reduceByKey()**
 - Combine values $[v_1, \dots, v_n]$ for key k into (k, v) where $v = \text{reduce}(v_1, \dots, v_n)$
 - Values for a key must be on the same partition/machine
- **Data shuffle** = re-distribute data across partitions/machines
- **Data shuffle is expensive** because of:
 - Data serialization (pickle)
 - Disk I/O (saving to disk)
 - Network I/O (copying across Executors)
 - Deserialization and memory allocation
- **Spark schedules general task graphs**
 - Automatic function pipelining
 - Data locality aware
 - Partitioning aware to avoid shuffles



Broadcast Variables

- **Challenge**

- Sending common variables to nodes with code can be costly
- Involves serialization, network transfer, deserialization
- Sending large, constant data repeatedly increases costs

- **Solution**

- Cache read-only variables on each node, to avoid repeated transfers

- **Example**

```
# `var` is large variable.  
var = list(range(1, int(1e6)))  
# Create a broadcast variable.  
broadcast_var = sc.broadcast(var)  
# Do not modify `var`, but use `broadcast_var.value` instead  
# of `var`.
```

Spark: Accumulators

- **Accumulator** is a shared variable used for aggregating values across tasks
 - Updated using associative and commutative operations (e.g., sum, max)
 - Efficient in parallel systems like MapReduce
- **Usage Example**
 - Accumulator initialized on the driver
 - Updated inside transformations (e.g., foreach) across workers
 - Value is collected back to the driver
 - Accumulator only guaranteed to update once per action
- **Example**

```
>>> accum = sc.accumulator(0)
>>> sc.parallelize([1, 2, 3, 4]).foreach(lambda x: accum.add(x))
>>> accum.value
```

10

Spark vs Hadoop MapReduce

- **Performance:** Spark faster
 - Processes data in-memory
 - Outperforms MapReduce, needs lots of memory
 - Hadoop MapReduce persists to disk after actions
- **Ease of use:** Spark easier to program
- **Data processing:** Spark more general

Gray Sort Competition

	Hadoop MR Record	Spark Record (2014)
Data Size	102.5 TB	100 TB
Elapsed Time	72 mins	23 mins
# Nodes	2100	206
# Cores	50400 physical	6592 virtualized
Cluster disk throughput	3150 GB/s	618 GB/s
Network	dedicated data center, 10Gbps	virtualized (EC2) 10Gbps network
Sort rate	1.42 TB/min	4.27 TB/min
Sort rate/node	0.67 GB/min	20.7 GB/min

- Sort benchmark, Daytona Gray: sort 100 TB of data (1 trillion records)
 - Spark-based System 3x faster with 1/10 the number of nodes
 - Speed up is ~30x
- [Ref](#)