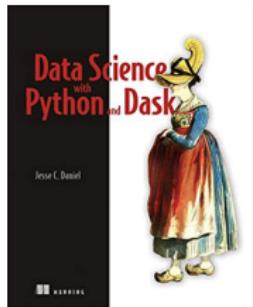


## 9.3: Python Dask

- **Instructor:** Dr. GP Saggese, [gsaggese@umd.edu](mailto:gsaggese@umd.edu)
- **Resources**
  - Web resources:
    - [Dask project](#)
    - [Dask examples](#)
  - Tutorial
    - [Dask tutorial](#)
    - [Dask advanced tutorial](#)
  - Class project
  - Mastery
    - Data science with Python and Dask, 2019



# Dataset Size Issues

---

- **Small datasets (< 1 GB)**
  - Fits into RAM
  - No disk paging needed
- **Medium dataset (< 1TB)**
  - Doesn't fit into RAM
  - Fits into local disk
    - Performance penalty with local disk
  - Need multiple CPU cores
    - Difficult to leverage parallelism with Python/Pandas
- **Large dataset (> 1TB)**
  - Doesn't fit into RAM
  - Doesn't fit into local disk
  - Need multiple servers
    - Python/Pandas not built for distributed datasets
    - Use frameworks for massive datasets
    - E.g., Hadoop, Spark, Dask, Ray



# Dataset Size Issues

---

Category	Size
Small datasets	< 1 GB
Medium datasets	< 1 TB
Large datasets	> 1 TB

- **The thresholds are fuzzy and changing over time**
  - Scale computer 10x to get 10x bigger datasets
- **Problem with scaling datasets**
  - Long run times
  - Rewriting code for different dataset sizes
  - Plan what and how to do efficiently
  - Cumbersome framework (Pandas easy, Hadoop difficult)

# Dask



- **Dask is written in Python**

- Scales Numpy, Pandas, sklearn
- Dask objects wrap library objects (e.g., Pandas DataFrame, numpy array)
- Parallel parts are “chunks” or “partitions”
  - Queued for work
  - Shipped between machines
  - Worked locally

- **Pros**


- Use familiar interfaces
- Write code optimized for parallelism
  - Dask handles heavy lifting

- **Scaling Dask is easy**

- Prototype on local machine, use cluster when needed
- No code refactoring needed
- No cluster-specific issues
  - E.g., resource management, data recovery, data movement
- Runs on multi-core
- Uses cluster managers
  - E.g., Yarn, Mesos, Kubernetes, AWS ECS



# Dask Layers



# Scaling Up vs Scaling Out

- **Scaling up**

- Replace equipment with larger, faster options
  - E.g., buy a larger pot, replace knife with food processor
- **Pros**
  - Better hardware, no code changes needed
- **Cons**
  - Exceed current machine capacity eventually
  - Cost: more powerful machines are expensive



- **Scaling out**

- Divide work between many workers in parallel
  - E.g., buy more pots and hire more cooks
- **Pros**
  - Task scheduler organizes computation, assigns workers to tasks
  - Cost-effective, no specialized hardware needed
- **Cons**
  - Write code to expose parallelism
  - Maintain cluster costs



# Dask: Computation

---

- **Lazy computations**

- Define transformations on data
- Define next computation without waiting
- Operate in chunks to avoid loading entire data in memory
- E.g.,
  - Split 2GB file into 32 64MB chunks
  - Operate on 8 chunks per server
  - Max memory use:  $512\text{MB} = (8 \times 64\text{MB})$
- Track object dimensions and data types
  - No code execution

- `compute()`

- Run computation (materialize)  
`missing_count_pct = missing_count.compute()`


- `persist()`

- Discard intermediate work to minimize memory
- Re-run graph for additional computation on intermediate nodes
- Keep intermediate result in memory
- Speed up large, complex DAGs for reuse

# Dask: Data Structures


- **Dask DataFrame**

- Implements Pandas DataFrame
- Tabular/relational data



- **Dask Array**

- Implements numpy ndarray
- Multidimensional array



- **Dask Bag**

- Coordinates Python lists of objects
- Parallelize computations on unstructured/semi-structured data

```
[1, 2, 3, 4, 5]  
[1, 2, 3] [4, 5]
```

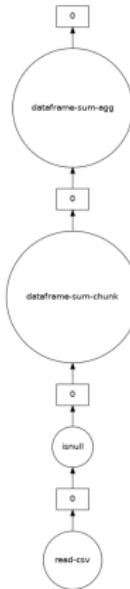
## Dask Reading Data

- Consider:

```
import dask.dataframe as dd
df = dd.read_csv('nyc-parking-tickets-2017.csv')
missing_values = df.isnull().sum()
missing_values
```

The screenshot shows a summary table for a DataFrame named 'df'. The columns are labeled: Column names, Summons number, Plate ID, Registration state, Plate type, Issue date, and Violation code. A callout points to the 'Summons number' column with the text 'npartitions=33'. Another callout points to the 'Plate ID' column with the text 'dtypes: int64 object'. A third callout points to the 'Registration state' column with the text 'object'. A fourth callout points to the 'Plate type' column with the text 'object'. A fifth callout points to the 'Issue date' column with the text 'object'. A sixth callout points to the 'Violations code' column with the text 'int64'. Below the table, a section titled 'Dask name: from-delayed, 99 tasks' is shown, with two arrows pointing upwards from the text 'Internal name of the Number of nodes in the underlying DAG' to the 'from-delayed' and '99 tasks' labels.

- `dask.dataframe.read_csv()`
    - Doesn't load data in memory
    - Infers column types
      - Samples data
      - Set data types
      - Use Parquet for data and types together
  - Partitions = independent data chunks
    - E.g., 33 partitions
    - Graph = 99 tasks
    - Each partition reads, splits data, initializes df object



# Low Level APIs: Delayed

- Handle computations that don't fit in native Dask data structures
  - E.g., Dask DataFrame
- In the example below there is parallelism that can be exploited

```
def inc(x):
    return x + 1


def double(x):
    return x * 2

def add(x, y):
    return x + y

data = [1, 2, 3, 4, 5]


output = []
for x in data:
    # (x + 1) + (x * 2) = 3x + 1
    a = inc(x)
    b = double(x)
    c = add(a, b)
    # 1 -> 4
    # 2 -> 7
    # 3 -> 10
    # 4 -> 13
    # 5 -> 16
    output.append(c)

# 4 + 7 + 10 + 13 + 16 = 20 + 20 + 10 = 50
total = sum(output)
print(total)
```



# Low Level APIs: Futures

- In parallel programming, a “future” encapsulates asynchronous execution, representing the eventual result
- Python `concurrent.futures`
  - High-level interface for asynchronous execution
  - Thread pool or Process pool (Executor interface)
- Task extends `concurrent.futures`
  - Express everything as futures
  - Specify blocking and non-blocking



```
def inc(x):
    return x + 1

def add(x, y):
    return x + y

a = client.submit(inc, 10)
b = client.submit(inc, 20)

>>> a
<Future: status: pending, key: inc-b8aaf26b99466a7a>

>>> a
<Future: status: finished, type: int, key: inc-b8aaf26b99466a7a>

>>> a.result()  # blocks until task completes and data arrives
```


11

# Different Types of Parallel Workload

- Break program in medium-size tasks of computation
  - E.g., a function call


**MapReduce**

Hadoop/Spark/Dask




**Embarrassingly Parallel**

Hadoop/Spark/Dask/Airflow/Prefect



**Full Task Scheduling**

Dask/Airflow/Prefect



# Encoding Task Graph

- Dask encodes tasks in terms of Python dicts and functions


```
def inc(i):
    return i + 1

def add(a, b):
    return a + b

x = 1
y = inc(x)
z = add(y, 10)
```



```
d = {'x': 1,
      'y': (inc, 'x'),
      'z': (add, 'y', 10)}
```



```
import dask.dataframe as dd

df = dd.read_csv('myfile.*.csv')
df = df + 100
df = df[df.name == 'Alice']
```



```
{
    # From the df.read_csv call
    ('read-csv', 0): (pandas.read_csv, 'myfile.0.csv'),
    ('read-csv', 1): (pandas.read_csv, 'myfile.1.csv'),
    ('read-csv', 2): (pandas.read_csv, 'myfile.2.csv'),
    ('read-csv', 3): (pandas.read_csv, 'myfile.3.csv'),

    # From the df + 100 call
    ('add', 0): (operator.add, ('read-csv', 0), 100),
    ('add', 1): (operator.add, ('read-csv', 1), 100),
    ('add', 2): (operator.add, ('read-csv', 2), 100),
    ('add', 3): (operator.add, ('read-csv', 3), 100),

    # From the df[df.name == 'Alice'] call
    ('filter', 0): (lambda part: part[part.name == 'Alice'], ('add', 0)),
    ('filter', 1): (lambda part: part[part.name == 'Alice'], ('add', 1)),
    ('filter', 2): (lambda part: part[part.name == 'Alice'], ('add', 2)),
    ('filter', 3): (lambda part: part[part.name == 'Alice'], ('add', 3)),
}
```

# Task Scheduling

- Data collections (Bags, Arrays, DataFrame) and operations create task graphs
  - Nodes: Python functions
  - Edges: Dependencies (output from one task used as input in another)
- Schedule task graphs for execution
  - Single-machine scheduler
    - Use local process or thread pool
    - Runs on a single machine
  - Distributed scheduler
    - Runs locally or across a cluster

## Collections

(create task graphs)




## Task Graph



## Schedulers

(execute task graphs)

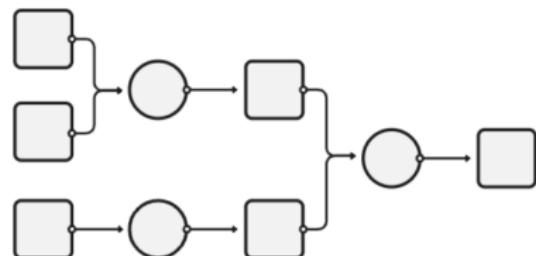


Single-machine  
(threads, processes,  
synchronous)

Distributed

# Task Scheduling

- **Dask task scheduler orchestrates work dynamically**
  - Not static scheduling like a relational DB
  - During computation, Dask dynamically assesses:
    - Completed tasks
    - Remaining tasks
    - Free resources (CPUs)
    - Data location
- **Dynamic approach handles various issues:**
  - Worker failure
    - Re-run tasks
  - Workers completing at different speeds due to:
    - Different computation
    - Different hardware
    - Varying server workloads
    - Slower data access
  - Network unreliability
    - Re-run or remove isolated nodes



# Dask vs Spark

---

- **Pros**

- Popular framework for large datasets
- In-memory alternative to MapReduce/Hadoop

- **Cons**

- Java library, supports Python via PySpark API
  - Python code runs on JVM
  - Debugging is difficult as execution is outside Python
- Different DataFrame API than Pandas
  - Learn “the Spark way”
  - May need to implement twice for exploratory analysis and production
- Optimized for MapReduce operations
- Difficult to set up and configure