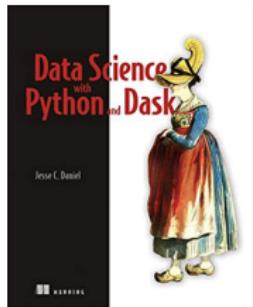


9.3: Python Dask

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu
- **Resources**
 - Web resources:
 - [Dask project](#)
 - [Dask examples](#)
 - Tutorial
 - [Dask tutorial](#)
 - [Dask advanced tutorial](#)
 - Class project
 - Mastery
 - Data science with Python and Dask, 2019



Dataset Size Issues

- **Small datasets (< 1 GB)**
 - Fits into RAM
 - No disk paging needed
- **Medium dataset (< 1TB)**
 - Doesn't fit into RAM
 - Fits into local disk
 - Performance penalty with local disk
 - Need multiple CPU cores
 - Difficult to leverage parallelism with Python/Pandas
- **Large dataset (> 1TB)**
 - Doesn't fit into RAM
 - Doesn't fit into local disk
 - Need multiple servers
 - Python/Pandas not built for distributed datasets
 - Use frameworks for massive datasets
 - E.g., Hadoop, Spark, Dask, Ray



Dataset Size Issues

Category	Size
Small datasets	< 1 GB
Medium datasets	< 1 TB
Large datasets	> 1 TB

- **The thresholds are fuzzy and changing over time**
 - Scale computer 10x to get 10x bigger datasets
- **Problem with scaling datasets**
 - Long run times
 - Rewriting code for different dataset sizes
 - Plan what and how to do efficiently
 - Cumbersome framework (Pandas easy, Hadoop difficult)

Dask



- **Dask is written in Python**

- Scales Numpy, Pandas, sklearn
- Dask objects wrap library objects (e.g., Pandas DataFrame, numpy array)
- Parallel parts are “chunks” or “partitions”
 - Queued for work
 - Shipped between machines
 - Worked locally

- **Pros**


- Use familiar interfaces
- Write code optimized for parallelism
 - Dask handles heavy lifting

- **Scaling Dask is easy**

- Prototype on local machine, use cluster when needed
- No code refactoring needed
- No cluster-specific issues
 - E.g., resource management, data recovery, data movement
- Runs on multi-core
- Uses cluster managers
 - E.g., Yarn, Mesos, Kubernetes, AWS ECS



Dask Layers



Scaling Up vs Scaling Out

- **Scaling up**

- Replace equipment with larger, faster options
 - E.g., buy a larger pot, replace knife with food processor
- **Pros**
 - Better hardware, no code changes needed
- **Cons**
 - Exceed current machine capacity eventually
 - Cost: more powerful machines are expensive



- **Scaling out**

- Divide work between many workers in parallel
 - E.g., buy more pots and hire more cooks
- **Pros**
 - Task scheduler organizes computation, assigns workers to tasks
 - Cost-effective, no specialized hardware needed
- **Cons**
 - Write code to expose parallelism
 - Maintain cluster costs



Dask: Computation

- **Lazy computations**

- Define transformations on data
- Define next computation without waiting
- Operate in chunks to avoid loading entire data in memory
- E.g.,
 - Split 2GB file into 32 64MB chunks
 - Operate on 8 chunks per server
 - Max memory use: $512\text{MB} = (8 \times 64\text{MB})$
- Track object dimensions and data types
 - No code execution

- `compute()`

- Run computation (materialize)
`missing_count_pct = missing_count.compute()`

- `persist()`

- Discard intermediate work to minimize memory
- Re-run graph for additional computation on intermediate nodes
- `persist()` keeps intermediate result in memory
- Speeds up large, complex DAGs for reuse

Dask: Data Structures


- **Dask DataFrame**

- Implements Pandas DataFrame
- Tabular/relational data



- **Dask Array**

- Implements numpy ndarray
- Multidimensional array



- **Dask Bag**

- Coordinates Python lists of objects
- Parallelize computations on unstructured/semi-structured data

[1, 2, 3, 4, 5]
[1, 2, 3] [4, 5]

Dask Reading Data

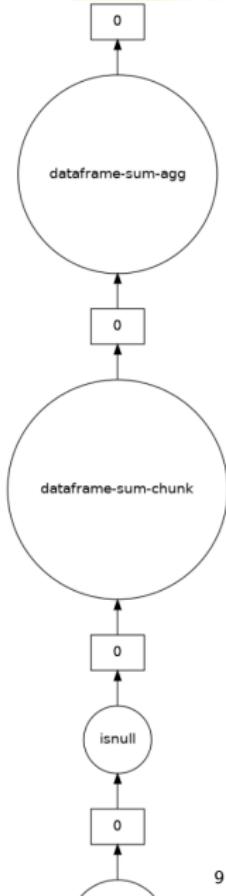
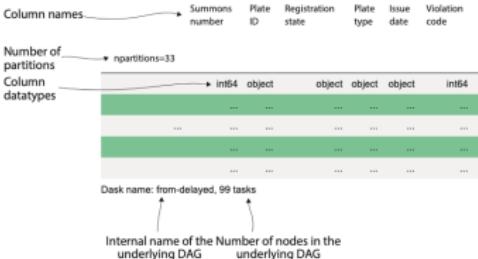
- 1

dask.dataframe.read_csv()

- Doesn't load the data in memory with
 - Tries to infer the types of the columns
 - By randomly sampling some data
 - Best to set the data types
 - Even better is to use Parquet since it stores data and types together

- Partitions = chunks of data that can be worked independently

- E.g., 33 partitions
 - Graph is composed of 99 tasks
 - Each partition reads data, splits data, initializes df object



Low Level APIs: Delayed

- Handle computations that don't fit in native Dask data structures (e.g., Dask DataFrame)
- In the example below there is parallelism that can be exploited

```
def inc(x):
    return x + 1


def double(x):
    return x * 2

def add(x, y):
    return x + y

data = [1, 2, 3, 4, 5]

output = []
for x in data:
    #  $(x + 1) + (x * 2) = 3x + 1$ 
    a = inc(x)
    b = double(x)
    c = add(a, b)
    # 1 -> 4
    # 2 -> 7
    # 3 -> 10
    # 4 -> 13
    # 5 -> 16
    output.append(c)

# 4 + 7 + 10 + 13 + 16 = 20 + 20 + 10 = 50
total = sum(output)
print(total)
```




Low Level APIs: Futures

- In parallel programming, a “future” encapsulates the asynchronous execution of a callable, representing the eventual result of the operation
- Futures is the most general way of specifying concurrency in Dask
 - Everything can be expressed in terms of futures
 - User can specify what's blocking and what's not blocking
-

Python **concurrent.futures**

- High-level interface for asynchronously executing callables
- Thread pool or Process



```
def inc(x):  
    return x + 1
```

```
def add(x, y):  
    return x + y
```

```
a = client.submit(inc, 10)  
b = client.submit(inc, 20)
```

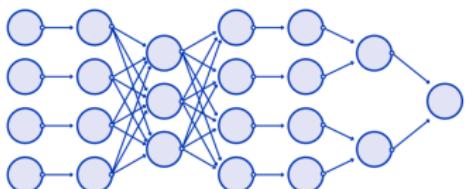
```
>>> a  
<Future: status: pending, key: inc-b8aaef26b99466a7a1980efafade6701d>  
>>> a  
<Future: status: finished, type: int, key: inc-b8aaef26b99466a7a1980efafade6701d>  
>>> a.result() # blocks until task completes and data arrives  
11
```

Different Types of Parallel Workload

- Break program in medium-size tasks of computation
 - E.g., a function call

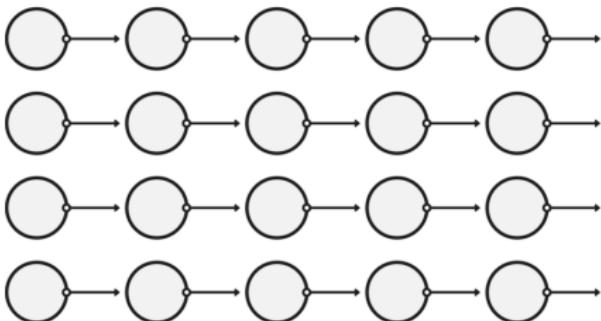
MapReduce

Hadoop/Spark/Dask




Embarrassingly Parallel

Hadoop/Spark/Dask/Airflow/Prefect



Full Task Scheduling

Dask/Airflow/Prefect



Encoding Task Graph

- Dask encodes tasks in terms of Python dicts and functions :: columns ::::

```
def inc(i):  
    return i + 1
```

```
def add(a, b):  
    return a + b
```

```
x = 1  
y = inc(x)  
z = add(y, 10)
```

Task Scheduling

- Data collections (Bags, Arrays, DataFrame) and their operations create task graphs
 - Nodes in the task graph are Python functions
 - Edges are dependencies (e.g., output from one task used as input in another task)
- Task graphs are scheduled for execution
- Single-machine scheduler
 - Use local process or thread pool
 - Simple but it can only run on a single machine
- Distributed scheduler
 - It can run locally or distributed across a cluster

Collections

(create task graphs)



Task Graph




Schedulers

(execute task graphs)

Dask Array

Dask DataFrame

Dask Bag

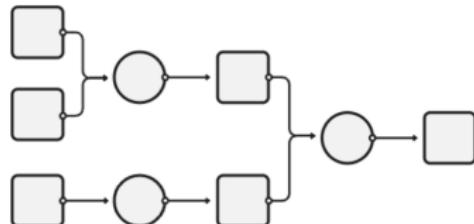


Single-machine
(threads, processes,
synchronous)

Distributed

Task Scheduling

- Dask task scheduler orchestrates the work dynamically
 - Not a static scheduling of operations like a relational DB
 - When the computation takes place, Dask dynamically assesses:
 - What tasks has been completed
 - What tasks is left to do
 - What resources (CPUs) are free
 - Where the data is located



- This dynamic approach handles a variety issues:

- Worker failure
 - Just re-run
- Workers completing work at different speeds

Dask vs Spark

- Spark has
- **Pros**
 - Popular framework for analyzing large datasets
 - In-memory alternative to MapReduce / Hadoop
- **Cons**
 - Spark is a Java library, supporting Python through PySpark API
 - Python code is executed on JVM through py4j
 - Difficult to debug since execution occurs outside Python
 - Different DataFrame API than Pandas
 - Learn how to do things “the Spark way”
 - You might need to implement things twice to go from exploratory analysis to large experiments / production
 - Optimized for MapReduce operations over a collection
 - Difficult to set-up and configure

Tutorial

Tutorial - From the official documentation

<https://docs.dask.org/en/stable/10-minutes-to-dask.html>