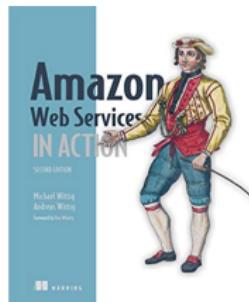




UMD DATA605 - Big Data Systems

11.2: AWS Overview

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu
- **References:**
 - Some basic info in the slides
 - Many tutorials on-line
 - Mastery
 - “Amazon Web Services in Action” 3rd Edition



Amazon Web Services (AWS)

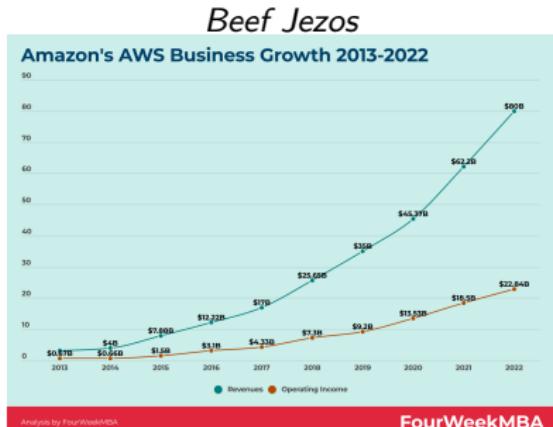
- AWS is a platform offering complete cloud solutions
 - Computing (e.g., EC2)
 - Storing (e.g., S3)
 - Networking
- Offers different levels of abstractions
 - IaaS, PaaS, SaaS
 - From virtual hardware to applications, e.g.,
 - Host websites
 - Run enterprise software
 - Run machine learning applications
- Control services in different ways
 - Web interface (console)
 - CLI: aws command
 - Language libraries, SDK (e.g., Python boto3)

AWS as Business

- Services charge pay-per-use pricing model
 - 500 new services and features annually
- Data centers distributed globally
 - US, Europe, Asia, South America
- Insanely profitable
 - \$91B/year in revenue (2023)
 - Grows 42% year-over-year
 - Controls 30% of cloud business



1998: "I sell books."
2017: "I sell whatever the f--- I want."



Types of Cloud Computing

- Cloud computing enables a **shared pool of configurable computing resources**
 - E.g., servers, storage, networks, applications, services
 - Accessible from anywhere
 - Convenient
 - On-demand
- Clouds can be:
 - *Public*: open to the general public (e.g., AWS)
 - *Private*: virtualize and share IT infrastructure within one organization (e.g., government)
 - *Hybrid*: mix of public and private clouds

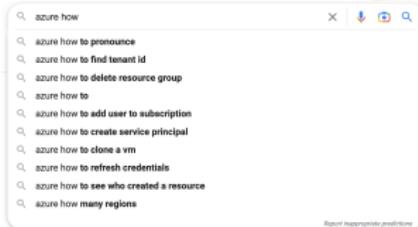
AWS vs Google Cloud vs Microsoft Azure

- **Similarities:**

- Worldwide infrastructure
- Provide IaaS (computing, networking, storage)
 - AWS EC2 / Google Compute Engine / Azure VMs
 - AWS S3 / Google Cloud Storage / Azure Blob Storage
- Pay-as-you-go pricing model

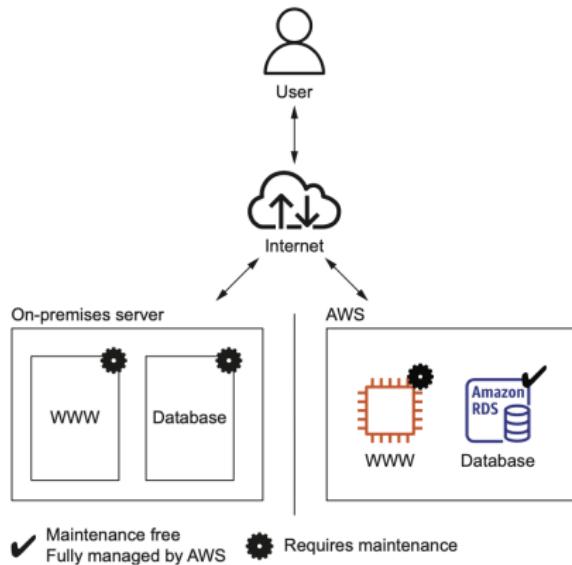
- **Differences:**

- AWS
 - Market leader, mature, powerful
 - Utilizes open source technologies
- Azure
 - Offers Microsoft stack in the cloud
- Google
 - Focuses on cloud-native applications



From On-premise to AWS: 1/3

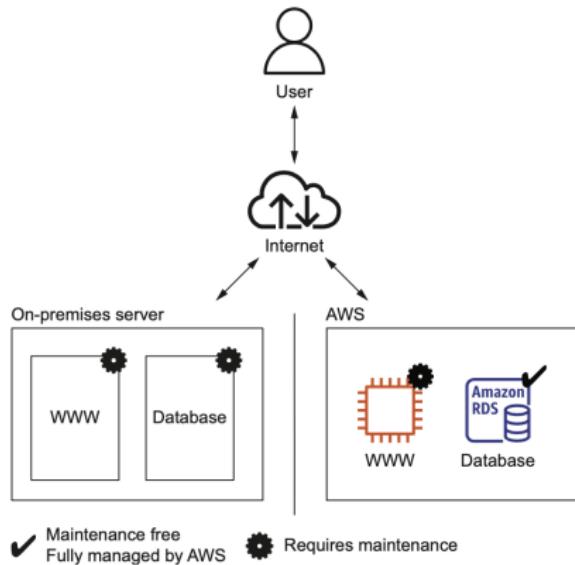
- Aka “Cloud-transformation”
- **Example:** move a medium-sized e-commerce site *from on-premise to the cloud*
- **No-cloud**
 - *Web-server*: handles customer requests
 - *DB*: stores product info and orders
 - *Static content* (e.g., JPEG images)
 - Delivered over CDN
 - Reduces load on services
 - *Dynamic content* (e.g., HTML pages)
 - E.g., products and prices
 - Delivered by web server



From On-premise to AWS: 2/3

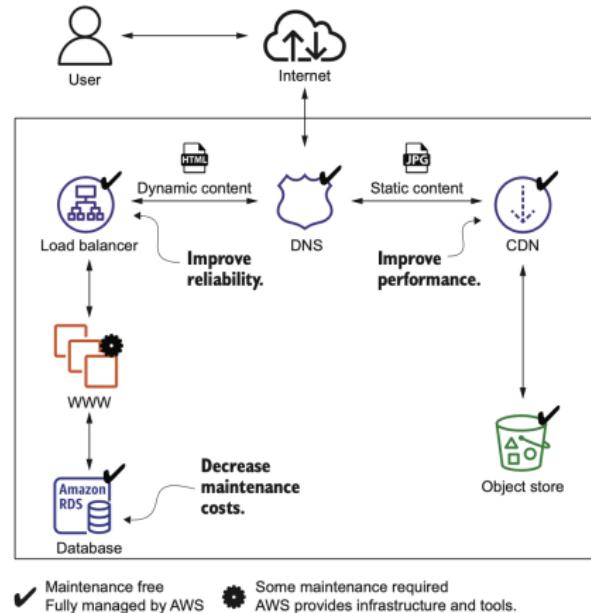
- **Move to cloud**

- Keep the same architecture
- Move components to cloud



From On-premise to AWS: 3/3

- Design for the cloud
 - DNS
 - Database
 - Object store (S3)
 - Managed solutions
 - Use multiple smaller virtual services with a load balancer to increase reliability



Low-Cost Processing

- Lots of batch jobs are not critical / run on a schedule
 - Analyze data daily
 - Generate reports from a database

1. **Allocate machines on demand**

- AWS bills VMs per minute
- Pay only when running jobs

2. **AWS Batch**

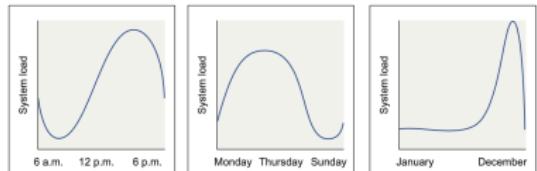
- Offers spare capacity at a discount
- Runs when capacity is available, saving 50%

3. **AWS Lambda**

- Serverless

Capacity Scaling

- **No need to plan for capacity**
 - Do not predict capacity
 - Schedule capacity on-the-fly
 - No concern for rack space, switches, power supplies
 - Add more VMs (1 to 1000) and storage (GB to PB) as needed
- **Handle seasonal traffic by scaling up/down**
 - Day vs night
 - Weekday vs weekend
 - Holidays
 - No test system needed when the team is off
- **Worldwide presence**
 - AWS has many data centers
 - Deploy applications close to customers



■ Analytics	■ Application integration	■ AR and VR
■ AWS cost management	■ Blockchain	■ Business applications
■ Compute	■ Containers	■ Customer enablement
■ Database	■ Developer tools	■ End-user computing
■ Frontend web and mobile	■ Game Development	■ Internet of Things
■ Machine learning	■ Management and governance	■ Media services
■ Migration and transfer	■ Networking and content delivery	■ Quantum technologies
■ Robotics	■ Satellite	■ Security, identity, and compliance
■ Storage		



Pay-per-use

- **AWS bill**

- Similar to an electric bill
- Services billed based on usage:
 - Hours of virtual server (rounded up)
 - Used storage in GB (allocated or real)
 - Data traffic in GB or requests
 - ...

- **Free tier**

- Use some AWS services free for 12 months after signing up
- Experiment with services
 - E.g., EC2, S3
- Exceed limits, start paying (without notice)
 - Set an alarm!

Service	January usage	February usage	February charge	Increase
Visits to website	100,000	500,000		
CDN	25 M requests + 25 GB traffic	125 M requests + 125 GB traffic	\$115.00	\$100.00
Static files	50 GB used storage	50 GB used storage	\$1.15	\$0.00
Load balancer	748 hours + 50 GB traffic	748 hours + 250 GB traffic	\$19.07	\$1.83
Web servers	1 virtual machine = 748 hours	4 virtual machines = 2,992 hours	\$200.46	\$150.35
Database (748 hours)	Small virtual machine + 20 GB storage	Large virtual machine + 20 GB storage	\$133.20	\$105.47
DNS	2 M requests	10 M requests	\$4.00	\$3.20
Total cost			\$472.88	\$360.65

Pay-per-use

- **Advantages**

- No upfront investments or commitment
- Lower project startup cost
- Easier to divide the system into smaller parts
 - One big server or two smaller ones cost the same
- Affordable fault tolerance/high performance
 - For a scalable workload “buy 1 server for 1000 hours” = “buy 1000 servers for 1 hour”



Interacting with AWS

• GUI (Management Console)

- Start interacting with services easily
- Set up cloud infrastructure for development and testing

• Command-line tool (CLI)

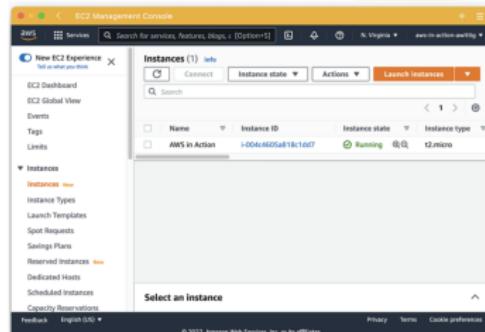
- Manage and access AWS services
- Automate recurring tasks

• SDKs

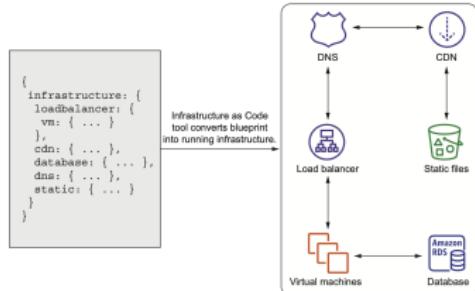
- Use libraries in any language to interact with AWS
- Integrate applications with AWS
- E.g., boto3 for Python

• Blueprints

- Describe your system with all services and dependencies
- Describe the system, not how to build it



```
andreas@stratus:~$ aws ec2 describe-instances
{
  "Reservations": [
    {
      "Group": "None",
      "Instances": [
        {
          "AmiLaunchIndex": 0,
          "ImageId": "ami-03291222c584d34f",
          "InstanceId": "i-004c6605a118c1a67",
          "InstanceType": "t3.nano",
          "LaunchTime": "2022-07-07T14:04:39+00:00",
          "Monitoring": {
            "State": "Disabled"
          },
          "Placement": {
            "AvailabilityZone": "us-east-1a",
            "Tenancy": "Default"
          }
        }
      ]
    }
  ]
}
```



Accounts and Users

- **Users**

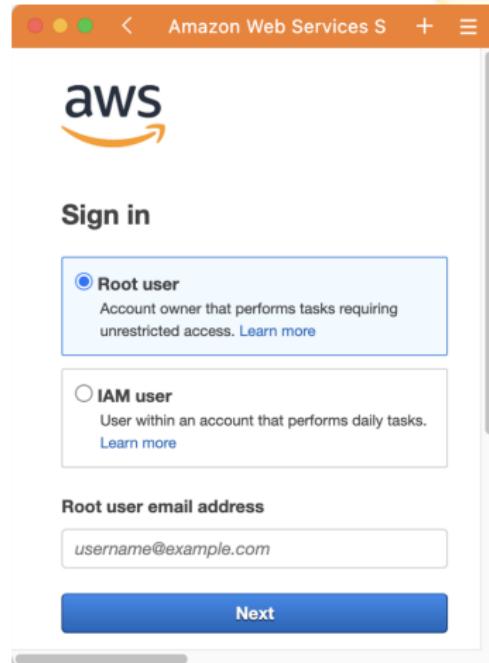
- One root AWS account
- Attach multiple users to an account
 - Different privileges
 - Isolate workloads

- **Be safe**

- Never use root account to develop
- Always use 2FA
- Avoid costly mistakes

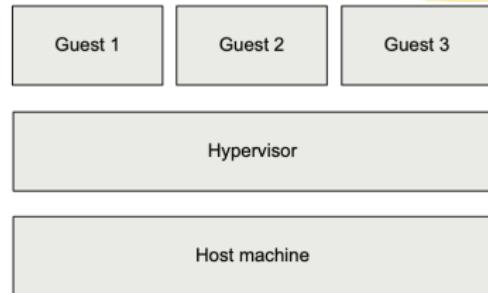
- **Key pair**

- Create a key pair to access a virtual server
- Public key (in AWS and on virtual servers)
- Private key is your secret
 - Don't lose it; can't retrieve it



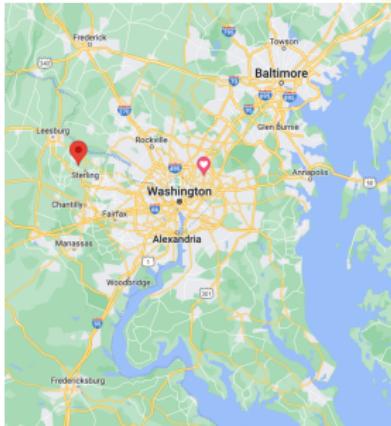
AWS VMs

- With virtualization
 - Multiple VMs run on the same hardware
 - Start and stop VMs on-demand
- Physical server**
 - Aka “host machine”, “bare metal”
 - Consists of CPUs, memory, networking interfaces, storage
- Hypervisor**
 - Software + CPU hardware
 - Isolates guests
 - Schedules hardware requests
- AWS**
 - Before Xen hypervisor (open-source)
 - Switched to AWS Nitro
 - Hardware-assisted virtualization
 - Performance close to bare metal
- Virtual servers**
 - Isolated on the same hardware



Starting an EC2 Instance

- **Select region**
 - E.g., us-east-1
 - 21155 Smith Switch Road, Ashburn, VA, USA



Starting an EC2 Instance

- **Select OS**

- Amazon Machine Image (AMI)
- Contains OS + pre-installed software
- Saves time installing OS, packages, software

- **Choose instance parameters**

- E.g., t2.micro
- More info later

- **Configure instance**

- Network, shutdown behavior, termination protection, monitoring

Select the dashboard.

Select N. Virginia region.

Start a virtual machine.

Instance type

t2.micro

Family t2 1 vCPU 1 GB Memory

On-Demand Linux pricing: 0.0176 USD per Hour

On-Demand Windows pricing: 0.0182 USD per Hour

Free tier eligible

Compare instance types

Application and OS Images (Amazon Machine Image)

Search our full catalog including 1000s of application and OS images

Recents

Quick Start

Select Amazon Linux.

Amazon Machine Image (AMI)

Amazon Linux 2 AMI (HVM) - Kernel 5.10, SSD Volume Type

ami-0c02fb55956c7d516

Virtualization type: HVM

ENA enabled: true

Root device type: ebs

Free tier eligible

Description

Amazon Linux 2 Kernel 5.10 AMI 2.0.20220316.0.x86_64 HVM gp2

Architecture

AMH ID

64-bit (x86)

ami-0c02fb55956c7d516

Use 64-bit (x86) architecture.

Choose Amazon Linux 2.

AWS Instance Type

- **Instance type** describes the computing power
- **Instance family**
 - T: cheap, baseline
 - M: general purpose
 - C: compute optimized
 - R: memory optimized
 - D: storage optimized for HDD
 - I: storage optimized for SSD
 - F: with FPGAs
 - P, G, CG: with GPUs
- [AWS EC2 list](#)

M7g	M4	M6g	M1	M1n	M1a	M5	M5n	M5m	M5a	M4	A1	T4g	T5
T3a	T2												
Instance	vCPU*		CPU Credits / hour		Mem (GiB)		Storage						Network Performance
t2.nano	1		3		0.5	EBS-Only							Low
t2.micro	1		6		1	EBS-Only							Low to Moderate
t2.small	1		12		2	EBS-Only							Low to Moderate
t2.medium	2		24		4	EBS-Only							Low to Moderate
t2.large	2		36		8	EBS-Only							Low to Moderate
t2.xlarge	4		54		16	EBS-Only							Moderate
t2.2xlarge	8		81		32	EBS-Only							Moderate

- E.g., t2.micro
 - t: instance family (small, cheap)
 - 2: generation (second)
 - micro: size (1 vCPU, 1GB memory)
 - 0.013 USD/hr
- E.g., m4.large
 - m: general purpose
 - large: size (2 vCPUs, 8GB memory)
 - 0.14 USD/hr

AWS Instance Type

- Price on AWS website is somehow unclear (duh really?)
 - Burst mode
 - vCPUs
 - Multi-tenancy
 - On-demand vs spot vs reserved vs pre-paid
- 642 types of EC2 machines (as of 2023)
 - Cheapest: 37 USD/yr
 - Most expensive: 1.91M USD/yr (500 CPUs and 24TB of memory)
- Alternative sites:
 - E.g., <https://instances.vantage.sh>

Starting an EC2 Instance

- Add storage
 - Volume size
 - Volume type (SSD or magnetic HDDs)
- Tag
- Configure firewall
 - How to access using SSH
 - Select key-pair
- How to monitor the instance
 - E.g. CloudWatch

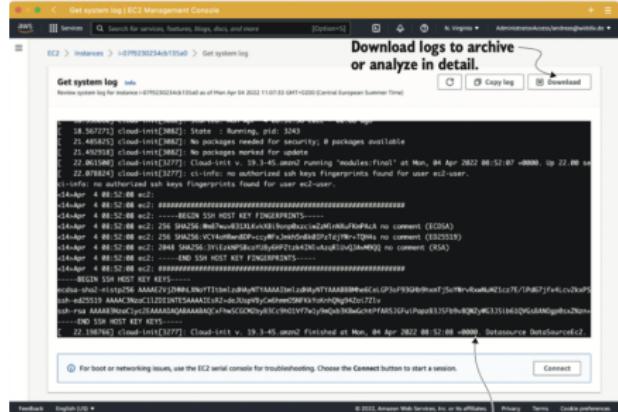
The screenshot shows two configuration steps in the AWS EC2 instance creation wizard.

Configure storage: Set to 1x 8 GiB gp2 Root volume. A note says "Configure 8 GB of storage for the root volume." A callout points to the volume type: "Use volume type gp2, which means your volume will store data on SSDs."

Network settings: Shows the selected network interface (vpc-a17392c7) and subnet. It indicates "Keep the default network." A callout points to the public IP section: "Ensure assigning a public IP is enabled." Another callout points to the security group creation: "Creates a new firewall configuration named launch-wizard-1".

Starting an EC2 Instance

- Start instance
- Find public IP



Get system log info
Review system log for instance i-07f9c30234a8155a6 as of Mon Apr 04 2022 11:07:33 GMT+0200 (Central European Summer Time)

Download logs to archive or analyze in detail.

Copy log Download

Get system log info

cloud-init v. 19.3-45.amzn2 running modules:final1 at Mon, 04 Apr 2022 08:52:07 +0000, ip 22.49.10.144

20.04.22/21:52 Cloud-Init [CRITICAL] Start: 1: running, pid: 3303

21.04.22/21:52 Cloud-Init [CRITICAL] No packages needed for security, 0 packages available

22.04.22/21:52 Cloud-Init [CRITICAL] Cloud-init v. 19.3-45.amzn2 running modules:final1 at Mon, 04 Apr 2022 08:52:07 +0000, ip 22.49.10.144

22.04.22/21:52 Cloud-Init [CRITICAL] ci-info: no authorized ssh keys fingerprints found for user ec2-user.

23.04.22/21:52 Cloud-Init [CRITICAL] ci-info: no authorized ssh keys fingerprints found for user ec2-user.

24.04.22/21:52 Cloud-Init [CRITICAL] 14-Apr 4:48:52.00 +0000 ec2: -----
----- BEGIN SSH HOST KEY FINGERPRINTS -----

24.04.22/21:52 Cloud-Init [CRITICAL] 14-Apr 4:48:52.00 +0000 ec2: 256 SHA256:VCh4t0WnB0DccyRwJxobD0h4BD7d1Nr-TDh4 no comment (RSA)
24.04.22/21:52 Cloud-Init [CRITICAL] 14-Apr 4:48:52.00 +0000 ec2: 256 SHA256:3V1zNP3kcr0jly46P7zxh4ZkNvAaUk1UvJaMwQ0 no comment (RSA)
24.04.22/21:52 Cloud-Init [CRITICAL] 14-Apr 4:48:52.00 +0000 ec2: ----- END SSH HOST KEY FINGERPRINTS -----

----- BEGIN SSH HOST KEY EXES -----

24.04.22/21:52 Cloud-Init [CRITICAL] 14-Apr 4:48:52.00 +0000 ec2: 256 AAAE2V3Y3H4LkW1TT1m1zJduUYtYMA23t1wduNtYTA888WwEc1Gf3cGP3t93Q4bRmuI5oWrxvewuR21cz7E/1p8d7f4Lcv2kxpS
24.04.22/21:52 Cloud-Init [CRITICAL] 14-Apr 4:48:52.00 +0000 ec2: ----- END SSH HOST KEY EXES -----
----- BEGIN SSH HOST KEYS -----

22.11.04/21:52 Cloud-Init [CRITICAL] 22.11.04/21:52 Cloud-Init [CRITICAL] Cloud-init v. 19.3-45.amzn2 Finished at Mon, 04 Apr 2022 08:52:08 +0000. Datasource DataSourceEc2.

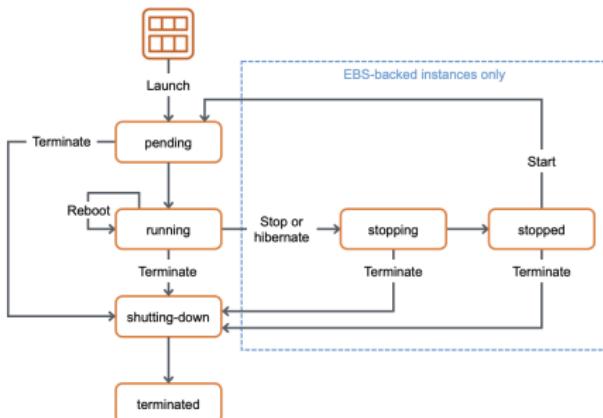
Feedback English (US) © 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookies preferences Connect

- Connect to the machine

```
> ssh -i $PATH/mykey.pem ubuntu@$PUBLIC_UP
> cat /proc/cpuinfo
> free -m
> sudo apt-get update
> sudo apt-get install ...
```

States of a VM

- **Start**
 - Start stopped VM
- **Stop**
 - Stopped VM not billed
 - Network HDD persist and incur charges
 - Local disk does not persist
 - VM can restart on a different host (different IP)
- **Reboot**
 - Network HDD persist
 - Local disk does not persist
 - Software remains installed
 - VM restarts on a different host
- **Terminate**
 - Delete: cannot restart
 - Network HDD persist
 - Local disk wiped out



Moving / Upgrading EC2 Instances

- **Scale up / down**

- Increase VM size for more computing power
- Stop VM
- Change instance type (e.g., m3.large)
- Start VM
- IP addresses change

- **AWS regions**

- Regions are collections of data centers
- Regions are independent
- No data transfer across regions
- Some services (e.g., IAM, CDN, DNS) are global

- **Why move across AWS regions**

- Proximity to users
- Compliance
 - Data storage and processing permissions
- Service availability
 - Some services unavailable in certain regions
- Redundancy
- Costs
 - Vary by region

Optimizing Costs

- **On-demand instances**
 - Maximum flexibility, no restrictions
 - Start and stop VMs anytime
 - Pay by hour
- **EC2 / Compute saving plans**
 - 1 yr vs 3 yrs
 - Commit to a certain number of hours
 - Payment options: all, partial, no upfront
 - Discount: up to 3x cheaper than on-demand
 - Useful for dev servers
- **Capacity reservation**
 - Access machines even in peak hours
- **Spot instances**
 - Bid for unused capacity
 - Price based on supply/demand
 - Discount: up to 10x cheaper than on-demand
 - Useful for asynchronous tasks



Programming the Infrastructure

- On AWS **everything can be controlled via an API**
 - Start a VM
 - Create 1TB of storage
 - Start Hadoop cluster
 - ...
- Jeff Bezos 2002 API mandate
 - An email worth \$100B/yr
- You can use:
 - AWS console
 - HTTP requests to the API
 - CLI (Command Line Interface)
 - Call AWS API from your terminal
 - SDK (Software Development Kit)
 - Call AWS API from your code
 - CloudFormation: templates describe infrastructure state, translated into API calls



Infrastructure-As-Code

- Use high-level programming language to control IT systems
- Apply software development principles to infrastructure
 - Code repository
 - Automated tests
 - Continuous integration
- DevOps (SRE in Google parlance)
 - Mix developers and operations in the same team
 - Use software to bring development and operations closer
 - Switch roles to experience each other's pain
 - Devs -> responsible for operational tasks (e.g., on-call)
 - Ops -> involved in software development, making system easier to operate
 - Foster communication and collaboration

Infrastructure-as-Code: Advantages

- **Save time**

- Reuse scripts or blueprints
- Automate regular tasks
- Copy-paste vs click-click-click

- **Fewer mistakes**

- Push-button flow

- **Consistency of actions**

- Multiple deployments per day

- **Deployment pipeline**

- Commit changes to source code
- Build application from source
- Automatic tests (e.g., integration tests)
- Build testing environment
- Run acceptance tests in isolation
- Propagate changes to production
- Monitor production

- **The script is detailed documentation**

- Explains what and how, not why (not a design document)

Create User Account

- Avoid using AWS root account for development
- Create a new user
 - IAM: Identity and Access Management
- Access key ID + secret access key
- Access control
 - Enable programmatic access
 - Enable console access
 - Limit user actions through policies

Select AdministratorAccess policy to grant full permissions.

Policy name	Type	Used as
<input checked="" type="checkbox"/> AdministratorAccess	Job function	Permissions policy (4)
<input type="checkbox"/> AdministratorAccess-Amplify	AWS managed	None
<input type="checkbox"/> AdministratorAccess-AWSLambdaBeanstalk	AWS managed	None
<input type="checkbox"/> AlexaForBusinessDeviceSetup	AWS managed	None
<input type="checkbox"/> AlexaForBusinessFullAccess	AWS managed	None

Set user details

You can add multiple users at once with the same access type and permissions. Learn more

User name* mycl
 Add another user

User name of the new user is mycl.

Select AWS access type

Select how these users will primarily access AWS. If you choose only programmatic access, it does NOT prevent users from accessing the console using an assumed role. Access keys and autogenerated passwords are provided in the last step. Learn more

Select AWS credential type*

Access key - Programmatic access

Enables an access key ID and secret access key for the AWS API, CLI, SDK, and other development tools.

Password - AWS Management Console access

Enables a password that allows users to sign-in to the AWS Management Console.

Check Programmatic access to generate an access key.

You haven't created a user at the moment.

Click to create a new user.

Bad idea!

AWS Command Line Interface (CLI)

- Provide unified interface to all AWS services
- Output is in JSON format

```
> apt-get install awscli
```

- Authenticate

```
> aws configure
```

- AWS access key ID
 - E.g., AKIAIR...D7ZCA
- AWS secret access key
 - E.g., SSKIng7jkA...enqBj7
- Default region name
 - E.g., us-east-1

- Execute command

```
> aws <service> <action> --key value
```

Software Development Kit (SDK)

- SDK is a library calling AWS API from your programming language
 - E.g., Python, Go, Ruby, C++, JavaScript
- **Pros:** handles
 - Authentication
 - Retry on error
 - HTTPS communication
 - XML / JSON de-/serialization
- **Cons**
 - Imperative approach
 - Deal with dependencies

```
import boto3
```

```
# Get the service resource.
```

```
dynamodb = boto3.resource('dynamodb')
```

```
# Create the DynamoDB table.
```

```
table = dynamodb.create_table(
```

```
    TableName='users',
```

```
    KeySchema=[
```

```
{
```

```
    'AttributeName': 'username',
```

```
    'KeyType': 'HASH'
```

```
},
```

```
{
```

```
    'AttributeName': 'last_name',
```

```
    'KeyType': 'RANGE'
```

```
}
```

```
],
```

```
AttributeDefinitions=[
```

```
{
```

```
    'AttributeName': 'username',
```

```
    'AttributeType': 'S'
```

```
},
```

```
{
```

```
    'AttributeName': 'last_name',
```

```
    'AttributeType': 'S'
```

```
},
```

```
],
```

```
ProvisionedThroughput={
```

```
    'ReadCapacityUnits': 5,
```

```
    'WriteCapacityUnits': 5
```

```
}
```

```
# Wait until the table exists.
```

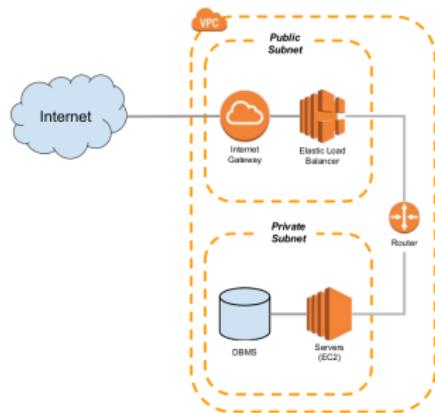
```
table.wait_until_exists()
```

AWS CloudFormation

- Use templates (e.g., JSON, YAML) to describe infrastructure
- Declarative vs imperative approach
 - Declare the system, not the steps to build it
- AWS Stacks processes CloudFormation templates
- Pros
 - Consistent infrastructure description
 - Handle dependencies
 - Customizable
 - Inject parameters to customize templates
 - Testable
 - Create infra from a template, test, shut down
 - Updatable
 - Update template, Stacks applies changes automatically
 - Serves as documentation
 - Use as code in source control

Securing Your System

- **Always install software updates**
 - Fix new security vulnerabilities daily
- **Restrict access to AWS account**
 - Use different AWS accounts for individuals and scripts
 - Apply “least privilege”: grant only necessary permissions
- **Restrict network traffic**
 - Open only essential ports
 - E.g., 80 for HTTP, 443 for HTTPS
 - Close all others
 - Encrypt traffic and data
- **Create a private network**
 - Use subnets not reachable from the Internet



AWS Shared-responsibility Principle

- AWS is responsible for:
 - Protecting the network through monitoring Internet access
 - E.g., preventing DDoS attacks
 - Ensuring the physical security of data centers
 - Decommissioning storage devices after the end of life
- You are responsible for:
 - Restricting access using IAM
 - Encrypting network traffic (e.g., HTTPS)
 - Configuring firewall for VPN
 - Encrypting data
 - Updating OS (kernel, libs) and software