



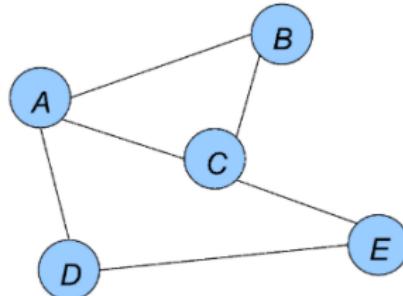
## UMD DATA605 - Big Data Systems

### 12.1: Graph Data Management

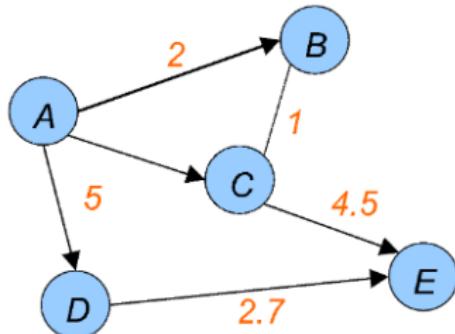
- **Instructor:** Dr. GP Saggese, [gsaggese@umd.edu](mailto:gsaggese@umd.edu)

# Graphs: Background

- A **graph** represents entities and their connections
  - Entities are *vertices* (or *nodes*)
  - Connections are *edges* (or *links*, *arcs*, *relationships*)
- **Applications in many fields**
  - Social networks
  - Biological networks
  - Information networks
  - Infrastructure networks
  - ...



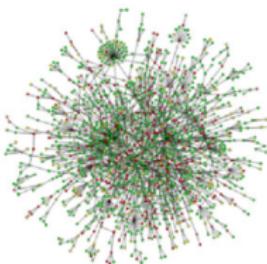
An *undirected, unweighted graph*



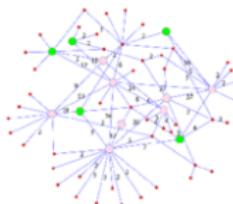
A *directed, edge-weighted graph*

# Graph Data Structures: Motivation

- Increasing volumes of graph data
- Increasing interest in querying and reasoning about graph data
- Sectors
  - Healthcare
  - Finance
  - Logistics
- Example applications
  - Fraud detection
  - Recommendation systems
  - Network analysis



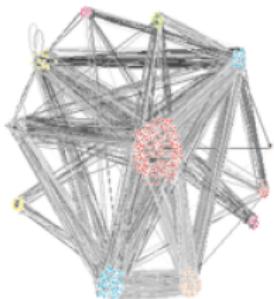
*Protein-protein  
interaction network*



*Supreme court  
citation network*



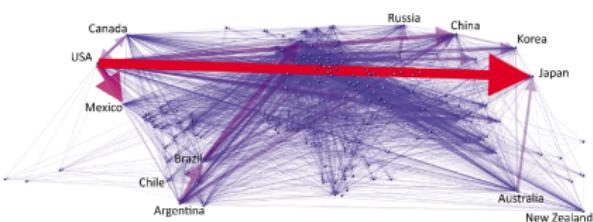
*Stock trading  
network*



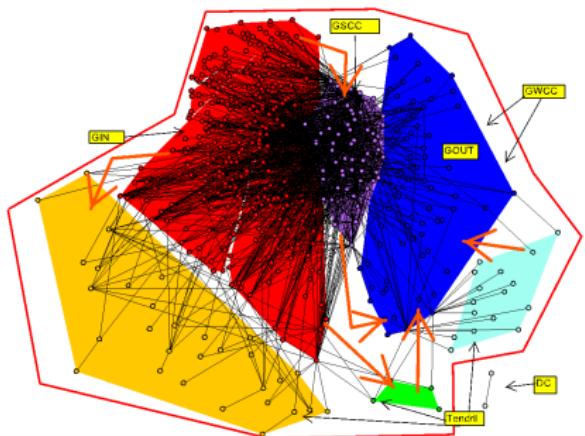
*Social networks*

# Graph Data Structures: Motivation

- **Traditional tools** (e.g., relational DBs, NoSQL DBs) struggle with:
  - Storing and querying graph data
    - Dedicated solutions: Neo4j
  - Processing graph-structured queries
    - Dedicated solutions: Google Pregel, Apache Giraph, Spark GraphX



*Global virtual trade network*



*Federal funds networks*

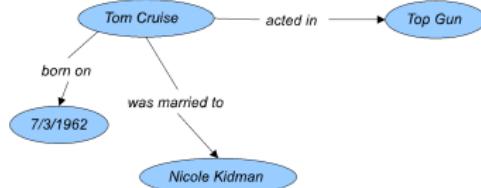
# Knowledge Graphs

---

- **Representation of knowledge in the form of graphs**
  - Capture entities, relationships, properties
  - Provide structured view of real-world information
- Represent using RDF or Property Graph models
  - E.g., Google Knowledge Graph, DBpedia, Wikidata
- **Applications**
  - Enable machine understanding of complex domains
  - Support semantic search, recommendation, analytics
  - Used in industries for data integration, knowledge discovery, AI applications
- **Ontologies**
  - Provide formal representation of knowledge
  - Promote interoperability across knowledge bases

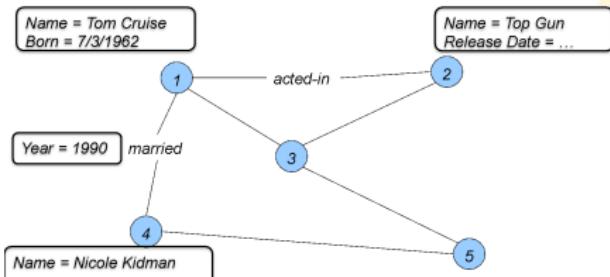
# Graph Data Models: RDF

- **Resource Description Framework**
  - RDF uses triples:  
subject-predicate-object
  - Connects “subject” and “object” through “predicate”
  - E.g., “TomCruise-acted-TopGun”
- **Used to represent knowledge bases**
  - Queried through SPARQL
- **Pros**
  - Standardization
    - W3C standard to model data
    - Subject and object can be URIs in semantic web
  - Interoperability
    - Merge RDF data stores
  - Extensibility
    - Add new nodes and relationships
    - Support ontologies



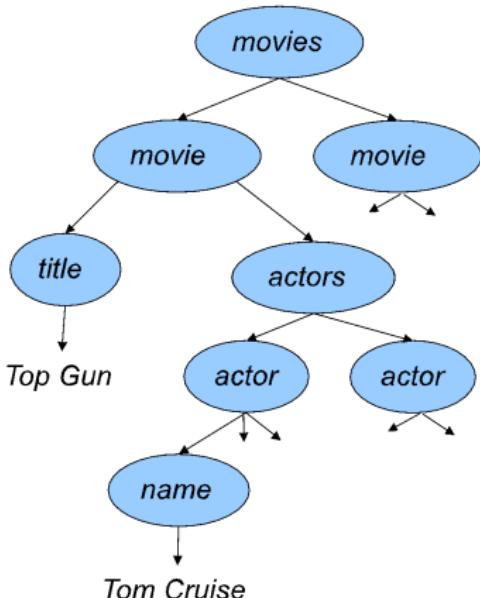
# Graph Data Models: Property Graph

- Directed graph with nodes and edges having key-values *properties*
  - Similar expressive power to RDFs
- **No universal standard**
  - Less “schema”
  - Harder to interoperate
- **Examples of query languages:**
  - Cypher for Neo4j
  - Gremlin for Apache TinkerPop



# Graph Data Models: XML

- Common data model for flexible data representation
- Directed labeled tree
- Popular for non-tabular data exchange



```
<movies>
  <movie>
    <title>Top Gun</title>
    <actors>
      <actor>
        <name>Tom Cruise</name>
        <born>7/3/1962</born>
      </actor>
      <actor>
        ...
      </actor>
    </actors>
  </movie>
</movies>
```

# Storing Graph Data

---

- **File systems**

- Very simple
- No support for transactions, ACID compliance
- Minimal functionality (e.g., must build the analysis/querying on top)

- **Relational databases**

- Mature technology
- All the good stuff (SQL, transactions, ACID compliance, toolchains)
- Minimal functionality for graph data

- **NoSQL key-value stores**

- Can handle very large datasets efficiently in a distributed fashion
- Minimal native functionality for graph data

- **Graph databases**

- Efficiently support complex queries/tasks (e.g., graph traversals)
- Not as mature as RDBMSs
- Often lack declarative language similar to SQL
  - You may need to write custom programs

# Graph Databases

---

- Many specialized graph database systems
  - E.g., Neo4j, Titan, OrientDB, AllegroGraph
- Key distinctions from relational databases
  - Manage and query graph-structured data
  - Store graph structure explicitly with pointers
    - Avoid joins, simplify graph traversals
    - Natural to write *queries* and *graph algorithms* (reachability, shortest paths)
  - Support graph query languages: SPARQL, Cypher, Gremlin
  - Rudimentary declarative interfaces
  - Applications often require programmatic interfaces
  - Provide programmatic API for arbitrary graph algorithms