



UMD DATA605 - Big Data Systems

12.1: Streaming and Real-time Analytics

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu

- ***Data Streams***
- Streaming Concepts
- Apache Streaming Zoo
- Processing Styles

Data Streams: Motivation

- Big Data is generated as a **continuous, unbounded stream**
- Applications generate data at **high velocity**
 - Financial transactions and market feeds
 - Sensor instrumentation, RFID, IoT telemetry
 - Network and system monitoring
 - Continuous media (video, audio)
- A **data stream** is a time-ordered sequence of events
 - Stream processing treats streams as first-class computational objects
- **Requirements**
 - Ingest and handle high-throughput event streams
 - Low-latency, near-real-time operations (e.g., time-series analytics)
 - Efficient dissemination of relevant subsets to consumers
 - Distributed processing to scale beyond a single machine

Data Streams: Examples

- Continuous queries
 - Any SQL query can be continuous
 - E.g., “*compute moving average over last hour every 10 mins*”
- Anomaly detection, pattern recognition
 - E.g., “*alert me when A occurs and then B within 10 mins*”
 - Correlate events from different streams
- Statistical tasks
 - E.g., de-noising measured readings
 - Build an online machine learning model
- Process multimedia data
 - E.g., online object detection, activity detection

Why Not Using Standard Solutions?

- **Example**


- “Report moving average of XYZ over last hour every 10 minutes”

- **Solution**

- Insert arriving items into a relational table
- Re-run query repeatedly

- **Problems**

- Re-executes full query, not leveraging incremental updates
- Many streaming computations are recursive
- Complex computations may not be easily expressed incrementally
- Real systems may run thousands of continuous queries




$$m_n = \frac{1}{n} \sum_{i=1}^n a_i$$

$$m_n = m_{n-1} + \frac{a_n - m_{n-1}}{n}$$

- Data Streams
- ***Streaming Concepts***
- Apache Streaming Zoo
- Processing Styles


Pub-Sub Systems: Motivation

- Modern distributed systems use **small, independent components**
 - E.g., serverless architectures, microservices (e.g., Uber)
 - Easier evolution, isolation, scalability
- **Publish-subscribe (pub-sub) systems**
 - Aka “message queues”, “message brokers”
 - Connect producers and consumers for event distribution
 - Topics cluster related messages
 - Typically provide lightweight dissemination rather than complex queries
 - Examples: AWS SQS, Kinesis, Kafka, RabbitMQ, Redis Streams, Celery, JBoss




Pub-Sub Systems: Architecture

- **Publishers**
 - Send messages or events
- **Subscribers**
 - Consume messages
- **Message broker**
 - Message broker routes event flow between publishers and subscribers, based on topics and subscriptions
- **Design parameters**
 - Event distribution model (topics, filters)
 - Push vs pull consumption
 - Subscriber interest patterns
 - Delivery guarantees
 - At-most-once
 - At-least-once
 - Exactly-once



Delivery Semantics: At-most once


- **At-most once:** *message may be lost, not redelivered*



- **Pros**
 - Small implementation overhead, high-performance
 - Easy to implement: "fire-and-forget"
- Works when occasional loss is acceptable
 - E.g., monitoring metrics of website

Delivery Semantics: At-least once


- **At-least once:** messages retried until acknowledged



- **Pros**
 - Ensures no loss but duplicates possible
- **Cons**
 - Requires idempotent operations or deduplication

Delivery Semantics: Exactly once


- **Exactly once:** each message processed once globally



- Most consumer-friendly but hardest to guarantee
 - Complicated by distributed coordination limits (e.g., "Two Generals' Problem")
- Used in financial and mission-critical systems
 - E.g., mission-critical systems (e.g., payment, trading, accounting)

Event vs Processing Time

- In both streaming and pub-sub architectures
 - **Event time**
 - Time when each record is generated
 - **Processing time**
 - Time when each record is received
 - Ingestion vs processing time: when events are received vs processed
- **Problems with events**
 - Events may arrive late or out of order
 - Determining how long to wait for stragglers is difficult
 - Systems set bounds on lateness
 - Extremely late data may be dropped or trigger re-computation




- Data Streams
- Streaming Concepts
- ***Apache Streaming Zoo***
- Processing Styles

Apache Streaming Zoo

- Many different streaming frameworks in the Apache family
 - E.g., Apex, Beam, Flink, Kafka, Spark, Storm, NiFi
 - Built simultaneously at different companies, then open-sourced
- **Different workloads**
 - Real-time analytics, continuous computation
 - Streaming ML, ETL pipelines
 - Messaging and log aggregation
- **Differences arise in**
 - Batch vs streaming orientation
 - Delivery semantics
 - Compute vs pub-sub roles
 - Throughput, latency, fault tolerance
 - API and language support

Apache Storm

- Open-source distributed real-time computation system
 - Acquired and open-sourced by Twitter
- **Horizontal scalability**: add machines to handle increasing data
- **Directed acyclic graph (DAG)**:
 - Spouts as data sources (as source nodes)
 - Bolts as processing units (as nodes)
 - Data streams (as edges)
- **Fault tolerance**:
 - At-least-once processing
 - Automatic task restarts
 - Workload redistribution
- **Suitable for**
 - Complex data processing workflows
 - With multiple stages and parallelism




Apache Kafka

- Open-source distributed streaming platform
 - Developed at LinkedIn, open-sourced in 2011




- **Core components**
 - Producers
 - Brokers
 - Consumers
 - Topics
 - Partitions
- **Delivery**: at-least-once, at-most-once, exactly-once
- High throughput, low latency
 - Persistent, replicated log storage
- *Kafka Connect* for integration with external systems
- *Kafka Streams* for native stream processing



Apache Flink




- Open-source, distributed data processing framework
- Distributed processing engine with strong support for stateful streaming
- Exactly-once semantics via checkpointing and robust state management
- Unified API for batch and streaming
- Rich windowing functions
- Runs on standalone clusters, YARN, Mesos, Kubernetes, and cloud



- Data Streams
- Streaming Concepts
- Apache Streaming Zoo
- ***Processing Styles***

Record-at-a-time Processing


- **Designed to handle infinite data streams**
 - Implemented in Apache Kafka
- **Distributed processing over multiple nodes**
 - Nodes organized in a DAG
 - Each node continuously:
 - Receives a single record
 - Processes the record immediately
 - Forwards the output to the next node



- **Pros**
 - Achieves extremely low latency
 - Example: sub-millisecond response times
- **Cons**
 - Poor fault tolerance
 - Requires extra nodes or redundant paths for failover
 - Sensitive to stragglers
 - Slow nodes can delay the entire pipeline

Micro-Batch Stream Processing

- Break continuous stream into **small batches** (e.g., 1-second windows)
 - Implemented in Spark Streaming (aka “DStreams”)



- **Pros**


- Recover from failures and stragglers with task scheduling
 - Schedule same task multiple times
- Deterministic tasks
 - Exactly-once processing
 - Consistent API: same semantics as RDDs
 - Fault-tolerance

- **Cons**

- Higher latency
 - E.g., seconds


Spark Structured Streaming

- Unified DataFrame/SQL-based model for both batch and streaming
- System manages state, faults, incremental computation, and late data
- Streaming table abstraction
 - Conceptually an unbounded table continuously appended with new rows
 - At time T, equivalent to a static DataFrame of all rows up to T



Incrementalization

- Framework identifies necessary state across micro-batches
- Uses DAG analysis to compute updated results from prior state
- Developers specify trigger conditions for updates
- Results updated incrementally as events arrive



Triggering Modes

- Indicate when to process newly available streaming data
- **Default**
 - Process micro-batch after previous completes
- **Trigger interval**
 - Specify fixed interval for each micro-batch
 - E.g., "every 10 minutes"
- **Once**
 - Wait for external trigger
 - E.g., "at end of day"
- **Continuous (experimental)**
 - Process data continuously
 - Not all operations available
 - Lower latency

Saving Modes

- Indicate when to save results and where
 - Each time result table updates, write to external file system
 - E.g., HDFS, AWS S3 or DB (e.g., MySQL, Cassandra)
- **Append mode**
 - Append new rows since last trigger
 - Use when existing rows don't change
- **Update mode**
 - Write updated rows since last trigger
 - Update in place
- **Complete mode**
 - Write entire updated result table
 - General but expensive

Spark Streaming “Hello world”

- `lines` is a `DataStreamReader`
 - Unbounded `DataFrame`
 - Set up reading but doesn't start reading
- `words` split data in words
- `counts` is a streaming `DataFrame`
 - Running word count
- `select()`, `filter()` are stateless transformations
- `count()` is stateful transformation
- Configuration
 - How to write processed output
 - Where to write (e.g., `console`)
 - How to write (e.g., `complete` for updated word counts)
 - When to trigger computation (e.g., every 1 second)
 - Where to save metadata for exactly-once guarantees, failure recovery
- `start()` processing (non-blocking)
 - `awaitTermination()` blocks until data is available

```
from pyspark.sql.functions import *
spark = SparkSession...
lines = (spark
    .readStream.format("socket")
    .option("host", "localhost")
    .option("port", 9999)
    .load())

words = lines.select(split(col("value"), "\\s").alias("word"))
counts = words.groupBy("word").count()
checkpointDir = "..."
streamingQuery = (counts
    .writeStream
    .format("console")
    .outputMode("complete")
    .trigger(processingTime="1 second")
    .option("checkpointLocation", checkpointDir)
    .start())
streamingQuery.awaitTermination()
```