



UMD DATA605 - Big Data Systems

12.1: Streaming and Real-time Analytics

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu

Motivation


- Big Data is generated as a **continuous, unbounded stream**
- Applications generate data at **high velocity**
 - Financial transactions and market feeds
 - Sensor instrumentation, RFID, IoT telemetry
 - Network and system monitoring
 - Continuous media (video, audio)
- A **data stream** is a time-ordered sequence of events
 - Stream processing treats streams as first-class computational objects
- **Requirements**
 - Ingest and handle high-throughput event streams
 - Low-latency, near-real-time operations (e.g., time-series analytics)
 - Efficient dissemination of relevant subsets to consumers
 - Distributed processing to scale beyond a single machine

Examples of Data Stream Tasks

- Continuous queries
 - Any SQL query can be continuous
 - E.g., “compute moving average over last hour every 10 mins”
- Anomaly detection, pattern recognition
 - E.g., “alert me when A occurs and then B within 10 mins”
 - Correlate events from different streams
- Statistical tasks
 - E.g., de-noising measured readings
 - Build an online machine learning model
- Process multimedia data
 - E.g., online object detection, activity detection

Why Not Using Standard Solutions?

- **Example**
 - “Report moving average of XYZ over last hour every 10 minutes”
- **Solution**
 - Insert arriving items into a relational table
 - Re-run query repeatedly
- **Problems**
 - Re-executes full query, not leveraging incremental updates
 - Many streaming computations are recursive
 - Complex computations may not be easily expressed incrementally
 - Real systems may run thousands of continuous queries




$$m_n = \frac{1}{n} \sum_{i=1}^n a_i$$

$$m_n = m_{n-1} + \frac{a_n - m_{n-1}}{n}$$


Pub-Sub Systems

- Modern distributed systems use **small, independent components**
 - E.g., serverless architectures, microservices (e.g., Uber)
 - Easier evolution, isolation, scalability
- **Publish-subscribe (pub-sub) systems**
 - Aka “message queues”, “message brokers”
 - Connect producers and consumers for event distribution
 - Topics cluster related messages
 - Typically provide lightweight dissemination rather than complex queries
 - Examples: AWS SQS, Kinesis, Kafka, RabbitMQ, Redis Streams, Celery, JBoss



Pub-Sub Systems: Architecture


- **Publishers**
 - Send messages or events
- **Subscribers**
 - Consume messages
- **Message broker**
 - Message broker routes event flow between publishers and subscribers, based on topics and subscriptions
- **Design parameters**
 - Event distribution model (topics, filters)
 - Push vs pull consumption
 - Subscriber interest patterns
 - Delivery guarantees
 - At-most-once
 - At-least-once
 - Exactly-once



Delivery Semantics: 1 / 3

- **At-most once**


- Message may be lost, not redelivered
- Small implementation overhead, high-performance
- Easy to implement: "fire-and-forget"
- Works when occasional loss is acceptable
- E.g., monitoring metrics of website



Delivery Semantics: 2 / 3

- **At-least once**


- Messages retried until acknowledged
- Ensures no loss but duplicates possible
- Requires idempotent operations or deduplication



Delivery Semantics: 2 / 3


- **Exactly once**

- Each message processed once globally
- Most consumer-friendly but hardest to guarantee
- Used in financial and mission-critical systems
- Complicated by distributed coordination limits (e.g., "Two Generals' Problem")
- E.g., mission-critical systems (e.g., payment, trading, accounting)



Event vs Processing Time

- In both streaming and pub-sub architectures
 - **Event time**
 - Time when each record is generated
 - **Processing time**
 - Time when each record is received
 - Ingestion vs processing time: when events are received vs processed
- **Problems with events**
 - Events may arrive late or out of order
 - Determining how long to wait for stragglers is difficult
 - Systems set bounds on lateness; extremely late data may be dropped or trigger recomputation



Apache Streaming Zoo

- Many different streaming frameworks in the Apache family
 - E.g., Apex, Beam, Flink, Kafka, Spark, Storm, NiFi
 - Built simultaneously at different companies, then open-sourced
- **Different workloads**
 - Real-time analytics, continuous computation
 - Streaming ML, ETL pipelines
 - Messaging and log aggregation
- **Differences arise in**
 - Batch vs streaming orientation
 - Delivery semantics
 - Compute vs pub-sub roles
 - Throughput, latency, fault tolerance
 - API and language support


Apache Storm

- Open-source distributed real-time computation system
 - Acquired and open-sourced by Twitter
- **Horizontal scalability**: add machines to handle increasing data
- **Fault tolerance**:
 - At-least-once** processing
 - Automatic task restarts
 - Workload redistribution
- **Directed acyclic graph (DAG)** with:
 - Spouts as data sources (as source nodes)
 - Bolts as processing units (as nodes)
 - Data streams (as edges)
- Suitable for complex data processing workflows with multiple stages and parallelism



Apache Kafka

- Open-source distributed streaming platform
 - Developed at LinkedIn, open-sourced in 2011
- Core components: producers, brokers, consumers, topics, partitions
- Persistent, replicated log storage
- High throughput, low latency
- Delivery: at-least-once, at-most-once, exactly-once
- Kafka Connect for integration with external systems
- Kafka Streams for native stream processing




Apache Flink

- Open-source, distributed data processing framework
- Focus on stateful computations over data streams
- Scalability: Horizontal scaling across large clusters
- Fault tolerance:
 - **Exactly-once** processing semantics, checkpointing, state management
- Batch processing support: **unified API for stream and batch processing**
- Flexible windowing: time-based, count-based, session windows
- Deployment options: standalone, YARN, Mesos, Kubernetes, cloud environments



Record-at-a-time Processing

- Implemented in Apache Kafka
- Goal: handle endless data stream
- **Multiple-node distributed processing engine**
 - Map computation on DAG of nodes
 - Each node continuously
 - Receives one record
 - Processes it
 - Forwards to next node
- **Pros**
 - Very low latencies
 - E.g., less than msec
- **Cons**
 - Inefficient node failure recovery
 - E.g., need failover resources/redundancy
 - Straggler nodes (slower than others)



Micro-Batch Stream Processing

- Aka DStreams
- Implemented in Spark Streaming
- Computation as a continuous series of batch jobs on small chunks of stream data


- E.g., 1 second
- Process each batch in the Spark cluster in a distributed manner

- **Pros**

- Recover from failures and stragglers using task scheduling
 - E.g., schedule same task multiple times
- Deterministic tasks
 - Exactly-once processing guarantees
 - Consistent API: same functional semantics as RDDs
 - Fault-tolerance

- **Cons**

- Higher latency
 - E.g., seconds




Spark Micro-Batch Processing: Cons

- Line between real-time and batch processing blurred
 - Application computing data hourly: stream or batch?
- No single API for batch and stream processing
 - Same abstractions (RDD) and operations
 - Rewrite code with different classes
- No support for event-time windows
 - Operations defined by processing time
 - No support for tardy data
- Spark replaced DStreams with Structured Streaming in v3
 - Supports micro-batch and continuous streaming
 - Closer batch vs streaming API


Spark Structured Streaming

- New approach used by Spark
- Goal: write stream processing as easy as writing batch pipelines
 - Single unified programming model
 - Use SQL or DataFrames on stream
- Handle automatically
 - Fault tolerance
 - Optimizations
 - Incremental computation
 - Tardy data
- **Data abstraction**
 - Batch applications: table (DataFrame) is the abstraction
 - Structured Streaming: table is an unbounded, continuously appended table
 - New record becomes a new row appended
 - At time T , it's like a static dataframe with data until T



Incrementalization

- Automatically detect state to maintain
 - Build DAG of computation
 - Express output of graph at time T in terms of graph at time $T-1$
 - Cache results
- Developers specify triggers to update results
- Incrementally update result with each record arrival



Triggering Modes

- Indicate when to process newly available streaming data
 - **Default**
 - Process micro-batch after previous completes
 - **Trigger interval**
 - Specify fixed interval for each micro-batch
 - E.g., "every 10 minutes"
 - **Once**
 - Wait for external trigger
 - E.g., "at end of day"
 - **Continuous (experimental)**
 - Process data continuously
 - Not all operations available
 - Lower latency

Saving Data

- Each time result table updates, write to external file system (e.g., HDFS, AWS S3) or DB (e.g., MySQL, Cassandra)
 - **Append mode**
 - Append new rows since last trigger
 - Use when existing rows don't change
 - **Update mode**
 - Write updated rows since last trigger
 - Update in place
 - **Complete mode**
 - Write entire updated result table
 - General but expensive

Spark Streaming “Hello world”

- `lines` looks like an RDD but it's a DataStreamReader
 - Unbounded DataFrame
 - Set up reading but doesn't start reading
- `words` split data in words
- `counts` is a streaming DataFrame
 - Running word count
- Stateless transformations don't require maintaining state
 - E.g., `select()`, `filter()`
- Stateful transformations
 - E.g., `count()`
- Define how to write processed output
 - Where to write (e.g., `console`)
 - How to write (e.g., `complete` for updated word counts)
- When to trigger computation
 - E.g., every 1 second
- Where to save metadata for:

```
from pyspark.sql.functions import *
spark = SparkSession...
lines = (spark
    .readStream.format("socket")
    .option("host", "localhost")
    .option("port", 9999)
    .load())

words = lines.select(split(col("value"), "\\s").alias("word"))
counts = words.groupBy("word").count()
checkpointDir = "..."
streamingQuery = (counts
    .writeStream
    .format("console")
    .outputMode("complete")
    .trigger(processingTime="1 second")
    .option("checkpointLocation", checkpointDir)
    .start())
streamingQuery.awaitTermination()
```

