




UMD DATA605 - Big Data Systems

12.2: Neo4j

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu

Neo4j

- Graph DB storing data as Property Graph
 - Nodes, edges hold data as key-value pairs
- Graph structure enables flexible schema
 - Focus is on relationships between values
- Two querying languages
 - Cypher
 - Gremlin
- GUI or REST API
- Full ACID-compliant transactions
- High-availability clustering
- Incremental backups
- Run in small application or large server clusters



Graph Data Model in Neo4j: Intuition

- **Nodes**

- Represent entities or objects
- Connected via *relationships*
- Have *properties* (key/value pairs)

- **Relationships**

- Represent (directional) connections between nodes
- Relationship types give semantic meaning to edges
- Multiple relationships per node
- Have *properties* (key/value pairs)

- **Properties**

- Store key–value information on nodes and relationships
 - Named values (key is a string)
- Indexed and constrained
- Composite indexes from multiple properties

- **Labels**

- Group nodes into sets with similar roles
- Nodes may have multiple labels
- Labels indexed for faster node retrieval
- Native label indexes optimized for performance

Why Cypher is Powerful

- Direct mapping between query and graph structure
 - Encourages thinking in relationships
 - Reduces impedance mismatch with graph data
- Scales naturally with connected data
- Enables expressive exploratory queries

Basic Cypher Pattern Matching

- Queries describe graph patterns to search for
 - Parentheses () represent nodes
 - Brackets [] represent relationships
 - Arrows -> or <- show relationship direction
- Example

(a)-[:FRIEND_OF]->(b)

MATCH Clause

- Used to find patterns in the graph
 - Similar to FROM ... WHERE in relational databases
 - Does not modify data
 - Can match multiple patterns in one query
- Example

```
MATCH (p:Person)-[:LIVES_IN]->(c:City)
```

Advanced Matching

- RETURN clause specifies what data to output
 - Can return nodes relationships or properties
 - Controls query result shape
 - Example
`RETURN p.name, c.name`
- Filtering with WHERE
 - Adds conditions to pattern matches
 - Works with properties labels and expressions
 - Often combined with MATCH
 - Example
`WHERE p.age > 30`
- Aggregation and Grouping
 - Uses functions like count, avg, max
 - Aggregation happens after MATCH
 - GROUP BY is implicit in RETURN
 - Example
`RETURN c.name, count(p)`

Creating Data with CREATE

- Used to add new nodes and relationships
- Pattern describes what should be created
- Executes exactly as written
- Example

CREATE

```
(a:Person {name:"Alice"})  
-[:KNOWS]->  
(b:Person {name:"Bob"})
```

Updating Graph Data

- SET modifies properties or labels
- REMOVE deletes properties or labels
- Allows incremental graph evolution
- Example

```
SET p.age = p.age + 1
```

Wine Suggestion Engine: Example 1/2

- **Create a wine suggestion engine**
 - Wines categorized by:
 - Varieties (e.g., Chardonnay, Pinot Noir)
 - Regions (e.g., Bordeaux, Napa, Tuscany)
 - Vintage (year grapes harvested)
 - Track articles describing wines by authors
 - Users track favorite wines
 - ...

Wine Suggestion Engine: Example 2/2

- **Relational approach**

- Create various tables


- wines: (id, name, year)
- wines_categories (wine_id, category_id)
- category table (id, name)
- wines_articles (wine_id, article_id)
- articles (id, publish_date, title, content)

- Relationships are

- produced
- reported_on
- grape_type

- **Problem with relational approach**

- There isn't much of a schema
- Lots of incomplete data
- An old saying in relational DB world: "*On a long enough timeline all fields become optional*"



Cypher Example

- **Graph DB approach:** provide values and structure only where necessary

```
CREATE (w:Wine
    {name: "Prancing Wolf",
     style: "ice wine",
     vintage: 2015})
```

```
CREATE (p:Publication
    {name: "Wine Expert Monthly"})
```

```
MATCH (p:Publication
    {name: "Wine Expert Monthly"}),
    (w:Wine {name: "Prancing Wolf",
    vintage: 2015})
CREATE (p)-[r:reported_on]->(w)
```

The screenshot shows a Cypher query being run in a graph database interface. The query is:

```
MATCH (a:Wine),(b:Publication) WHERE a.name = 'Prancing Wolf' AND b.name = 'Wine Expert Month...'.
```

The results pane shows two nodes found:

- Publication(1)
- Wine(1)


Below the results, there are tabs for Graph, Rows, and Text, with the Graph tab selected. A sidebar on the left contains icons for Graph, Rows, Text, and a search function.

Cypher Example

```
MATCH (p:Publication {name: "Wine Expert Monthly"}),  
      (w:Wine {name: "Prancing Wolf"})  
CREATE (p)-[r:reported_on {rating: 2}]->(w)
```

```
CREATE (g:GrapeType {name: "Riesling"})
```

```
MATCH (w:Wine {name: "Prancing Wolf"}),  
      (g:GrapeType {name: "Riesling"})  
CREATE (w)-[r:grape_type]->(g)
```



Cypher Example

```
CREATE (wr:Winery {name: "Prancing Wolf Winery"})
MATCH (w:Wine {name: "Prancing Wolf"}),
      (wr:Winery {name: "Prancing Wolf Winery"})
CREATE (wr)-[r:produced]->(w)
CREATE (w:Wine
       {name:"Prancing Wolf", style: "Kabinett", vintage: 2002})
CREATE (w:Wine
       {name: "Prancing Wolf", style: "Spätlese", vintage: 2010})
MATCH (wr:Winery
       {name: "Prancing Wolf Winery"}),(w:Wine {name: "Prancing Wolf"})
CREATE (wr)-[r:produced]->(w)
MATCH (w:Wine), (g:GrapeType {name: "Riesling"})
CREATE (w)-[r:grape_type]->(g)
```



Cypher Example

- Add a social component to the wine graph
 - People preference for wine
 - Relationships with one another
- The changes were made “superimposing” new relationships without changing the previous data ::::columns ::::{.column width=50%}

```
CREATE (p:Person {name: "Alice"})
```

```
MATCH (p:Person {name: "Alice"}),  
      (w:Wine {name: "Prancing Wolf",  
                style: "ice wine"})  
CREATE (p)-[r:likes]->(w)
```

```
CREATE (p:Person {name: "Patty"})
```

```
MATCH (p1:Person {name: "Patty"}),  
      (p2:Person {name: "Tom"})  
CREATE (p1)-[r:friends]->(p2)
```

:::: ::::{.column width=45%}



Cypher: Query Structure


```
MATCH [Nodes and relationships]
WHERE [Boolean filter statement]
RETURN [DISTINCT] [statements [AS alias]]
ORDER BY [Properties] [ASC or DESC]
SKIP [Number] LIMIT [Number]
```

Cypher: Query Example

```
MATCH (p:Person  
{name: "Alice"})-->(n)  
RETURN n;
```

```
MATCH (p:Person  
{name: "Alice"})-->(other: Person)  
RETURN other.name;
```

```
MATCH (fof:Person)-[:friends]-(f:Person)-[:friends]-(p:Person {name: "Alice"})  
RETURN fof.name;
```



Matching nodes and relationships

- Nodes
 - (a), (), (:Ntype), (a:Ntype),
 - (a { prop:'value' }) ,
 - (a:Ntype { prop:'value' })
- Relationships
 - (a)-(b)
 - (a)->(b), (a)<-(b),
 - (a)->(), (a)-[r]->(b),
 - (a)-[:Rtype]->(b), (a)-[:R1|R2]->(b),
 - (a)-[r:Rtype]->(b)
- May have more than 2 nodes
 - (a)->(b)<-(c), (a)->(b)->(c)
- Path
 - $p = (a) \rightarrow (b)$

More options

- Relationship distance:
 - (a)-[:Rtype*2]->(b): 2 hops of type Rtype
 - (a)-[:Rtype*]->(b): any number of hops of type Rtype
 - (a)-[:Rtype*2..10]->(b): 2-10 hops of Rtype
 - (a)-[:Rtype*..10]->(b): 1-10 hops of Rtype
 - (a)-[:Rtype*2..]->(b): at least 2 hops of Rtype
- Could be used also as:
 - (a)-[r*2]->(b) r gets a sequence of relationships
 - (a)-[*{prop:val}]->(b)

Operators

- Mathematical
 - +, -, *, /, %, ^ (power, not XOR)
- Comparison
 - =, <, >, <>, >=, <=, =~ (Regex), IS NULL, IS NOT NULL
- Boolean
 - AND, OR, XOR, NOT
- String
 - Concatenation through +
- Collection
 - Concatenation through +
 - IN to check if an element exists in a collection

More WHERE options

- WHERE others.name IN ['Andres', 'Peter']
- WHERE user.age IN range (18,30)
- WHERE n.name =~ 'Tob.*'
- WHERE n.name =~ '(?i)ANDR.*' - (case insensitive)
- WHERE (tobias)->()
- WHERE NOT (tobias)->()
- WHERE has(b.name)
- WHERE b.name? = 'Bob' (Returns all nodes where name = 'Bob' plus all nodes without a name property)

Functions

- On paths:
 - MATCH shortestPath((a)-[*]-(b))
 - MATCH allShortestPath((a)-[*]-(b))
 - Length(path) – The path length or 0 if not exists.
 - RETURN relationships(p) - Returns all relationships in a path.
- On collections:
 - RETURN a.array, filter(x IN a.array WHERE length(x)= 3) FILTER - returns the elements in a collection that comply to a predicate.
 - WHERE ANY (x IN a.array WHERE x = “one”) – at least one
 - WHERE ALL (x IN nodes(p) WHERE x.age > 30) – all elements
 - WHERE SINGLE (x IN nodes(p) WHERE var.eyes = “blue”) – Only one
- nodes(p) – nodes of the path p

With

- Manipulate result sequence before passing to following query parts
- Usage of WITH:
 - Limit entries passed to other MATCH clauses
 - Introduce aggregates for predicates in WHERE
 - Separate reading from updating the graph. Each query part must be read-only or write-only

Data access is programmatic

- REST API
- Through the Java APIs
 - JVM languages have bindings to the same APIs
 - JRuby, Jython, Clojure, Scala...
- Managing nodes and relationships
- Indexing
- Traversing
- Path finding
- Pattern matching