



UMD DATA605 - Big Data Systems

12.3: Graph Data Processing

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu

Options for Processing Graph Data

- Write your own programs
 - Extract relevant data, construct in-memory graph
 - Different storage options help to varying degrees
- Write queries in a declarative language
 - Suitable for some graph queries/tasks
 - E.g., Cypher for Neo4j
- Use a general-purpose distributed programming framework
 - E.g., Hadoop or Spark
 - Difficult for many graph analysis tasks
- Use a graph-specific programming framework
 - Simplifies writing graph analysis tasks, scales to large volumes
 - E.g., Giraph or GraphX

Option 2: Declarative Interfaces

- No consensus on declarative, high-level languages for querying or analysis
 - Variety in query/analysis tasks
 - Hard to find and exploit commonalities
- Limited, useful solutions:
 - XQuery for XML
 - Limited to tree-structured data
 - SPARQL for RDF
 - Standardized query language, limited functionality
 - Cypher by Neo4j
 - Datalog-based frameworks for analysis tasks
 - Many prototypes, task-specific

Option 3: MapReduce

- Popular option for processing large datasets
 - Hadoop or Spark
- Key advantages:
 - Scalability without scheduling, distributed execution, fault tolerance concerns
 - Simple programming framework
- Disadvantages:
 - Difficult for graph analysis tasks
 - Each traversal requires a new map-reduce phase
 - Hadoop not ideal for many phases, Spark is better
- Much work on graph analysis tasks using MapReduce

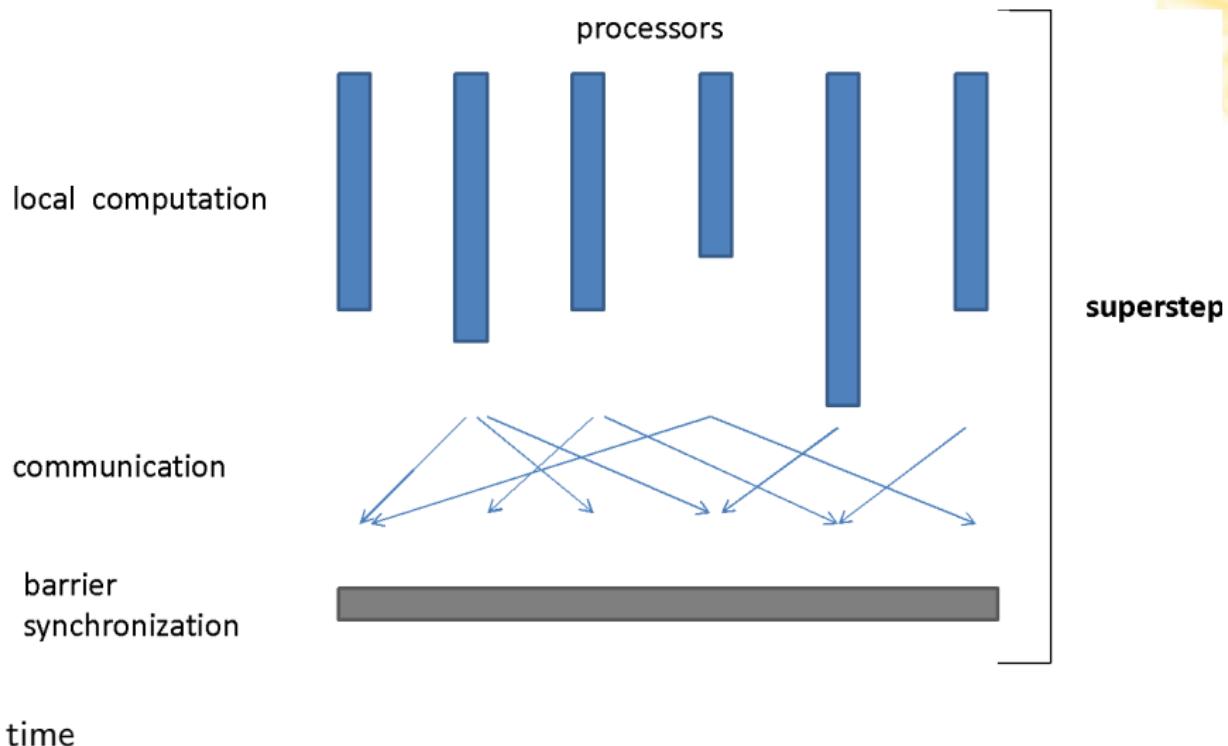
Option 3: MapReduce

- Disadvantages:
 - Difficult for graph analysis
 - Each traversal requires a new map-reduce phase
 - Each job executes N times
 - Hadoop not ideal for many phases (even with YARN)
 - Inefficient – redundant work
 - Mappers send PR values and graph structure
 - In PageRank: repeated reading and parsing each iteration
 - Extensive I/O at input, shuffle/sort, output

Option 4: Graph Programming Frameworks


- Frameworks (analogous to MapReduce) for analyzing large graph data
 - Address MapReduce limitations
 - Most are *vertex-centric*
 - Programs from a vertex perspective
 - Based on message passing between nodes
- Pregel: original framework by Google
 - Based on “Bulk Synchronous Parallel” (BSP) model
- Giraph: open-source Pregel on Hadoop
- GraphLab: asynchronous execution
- GraphX: built on Spark

Bulk Synchronous Parallel (BSP)



Vertex-centric BSP

- Each vertex has an id, value, list of adjacent vertex ids, and edge values
- Each vertex invoked in each superstep, recomputes value, sends messages to other vertices, delivered over superstep barriers
- Advanced features: termination votes, combiners, aggregators, topology mutations



Think like a vertex

- I know my local state
- I know my neighbours
- I can send messages to vertices
- I can declare that I am done
- I can mutate graph topology

Option 4: Pregel


- Programmers write one program: `compute()`
- Typical structure of `compute()`:
 - *Inputs*: current values of the node
 - *Inputs*: messages from neighboring nodes
 - Modify current values (if desired)
 - *Outputs*: send messages to neighbors
- Execution framework:
 - Execute `compute()` for all nodes in parallel
 - Synchronize (wait for all messages)
 - Repeat

Apache Giraph


- Pregel is proprietary, but:
 - **Apache Giraph**: open source implementation
 - Runs on standard **Hadoop** infrastructure
 - Computation executed in memory
 - Can be a job in a pipeline (**MapReduce, Hive**)
 - Uses **Apache ZooKeeper** for synchronization
 - Graph partition via hashing
 - Fault tolerance via checkpointing

Plays well with Hadoop

Giraph



Giraph Execution



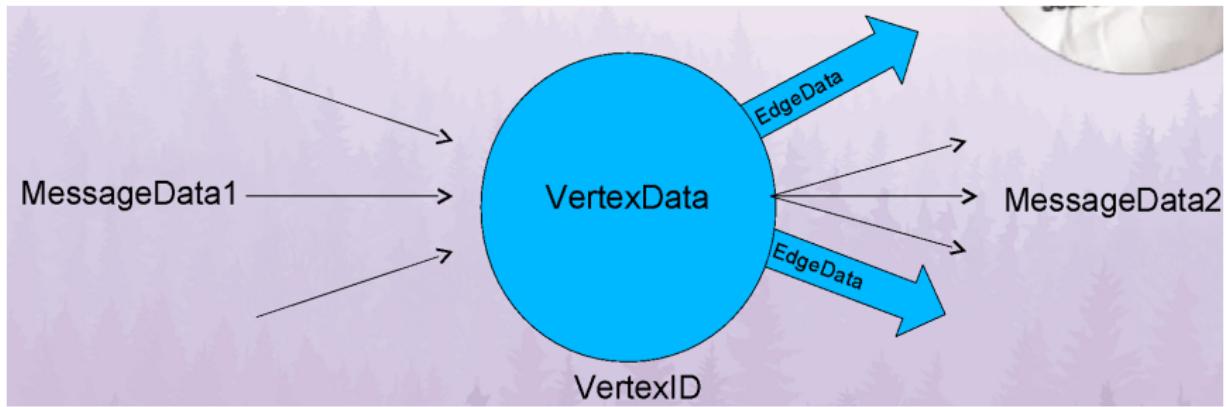
Which part is doing what?

- **ZooKeeper**: responsible for computation state
 - partition/worker mapping
 - global state: #superstep
 - checkpoint paths, aggregator values, statistics
- **Master**: responsible for coordination
 - assigns partitions to workers
 - coordinates synchronization
 - requests checkpoints
 - aggregates aggregator values
 - collects health statuses
- **Worker**: responsible for vertices
 - invokes active vertices compute() function
 - sends, receives, assigns messages
 - computes local aggregation values

What do you have to implement?

- Your algorithm as a **Vertex**
 - Subclass existing implementations: `BasicVertex`, `MutableVertex`, `EdgeListVertex`, `HashMapVertex`, `LongDoubleFloatDoubleVertex`
- A `VertexInputFormat` to read your graph
 - e.g., from a text file with adjacency lists like ...
- A `VertexOutputFormat` to write back the result
 - e.g.,

A vertex view



Designed for iterations

- Stateful (in-memory)
 - Keep all data in memory if possible
- Only intermediate values (messages) sent
 - Communicate with other vertices
- Hits disk at input, output, checkpoint
- Can go out-of-core
 - If data doesn't fit into memory

Graph modeling in Giraph

- BasicComputation< I extends WritableComparable, // VertexID – vertex ref V extends Writable, // VertexData – a vertex datum E extends Writable, // EdgeData – an edge label M extends Writable> // MessageData-- message payload

Giraph “Hello World”

```
public class GiraphHelloWorld extends
BasicComputation<IntWritable, IntWritable, NullWritable, NullWritabl

public void compute(Vertex<IntWritable, IntWritable, NullWritabl
    System.out.println("Hello world from the: " + vertex.getId() +
        System.out.println(" " + e.getTargetVertexId());    }
    System.out.println("");
    vertex.voteToHalt();
}
}
```

Example: Ping neighbors

```
public void compute(Vertex<Text, DoubleWritable, DoubleWritable> v  
    if (getSuperstep() == 0) {           sendMessageToAllEdges(vertex,  
} else {      for (Text m : ms) {  
        if (vertex.getEdgeValue(m) == null) {  
    }     vertex.voteToHalt();  
}
```

Giraph PageRank Example

```
public class PageRankComputation extends BasicComputation<Int  
    //Number of supersteps  
    public static final String SUPERSTEP_COUNT = "girap
```

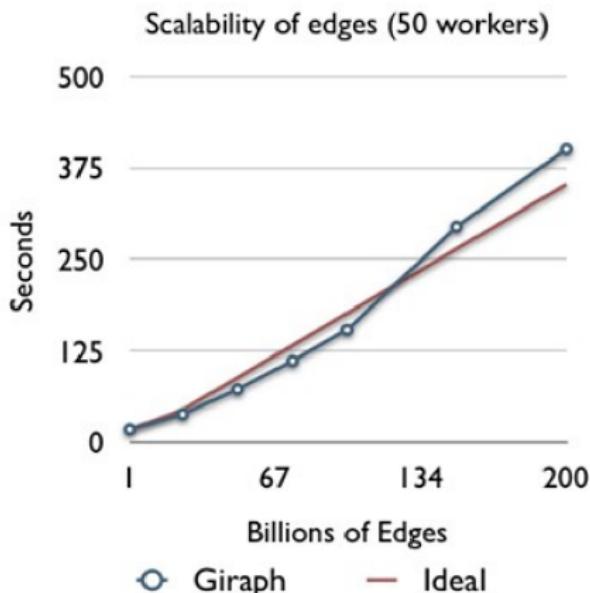
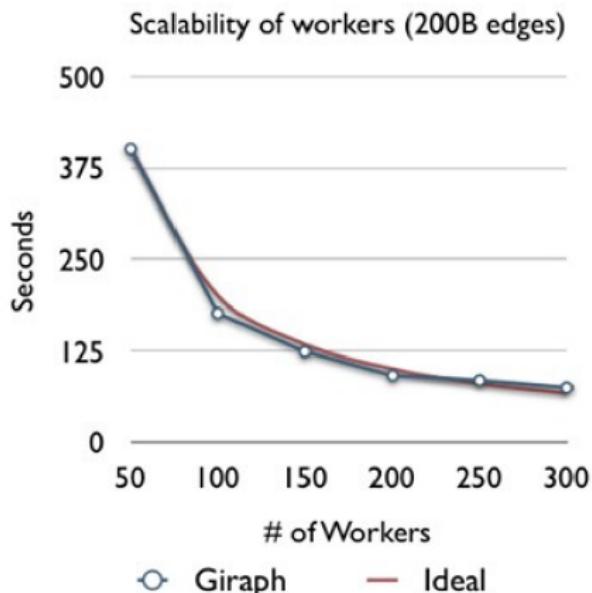
Giraph PageRank Example

```
public void compute(Vertex<IntWritable, FloatWritable, NullWritable> vertex, MapReduceInputContext<IntWritable, FloatWritable, NullWritable> context) throws IOException, InterruptedException {
    if (getSuperstep() >= 1) {
        float sum = 0;
        for (FloatWritable message : messages) {
            sum += message.get();
        }
        vertex.getValue().set((0.15f / getTotalNumVertices()) +
        sum);
        if (getSuperstep() < getConf().getInt(SUPERSTEP_COUNT, 0)) {
            sendMessageToAllEdges(vertex,
                new FloatWritable(vertex.getValue().get() /
                messages.size()));
        } else {
            vertex.voteToHalt();
        }
    }
}
```

Additional functionality

- Combiners
 - To minimize messages
- Aggregators
 - global aggregations across vertices
- MasterCompute
 - computation executed on master
- WorkerContext
 - executed per worker task
- PartitionContext
 - executed per partition

Giraph scales




<https://www.facebook.com/notes/facebook-engineering/scaling-apache-giraph-to->


Graphx

GraphX Motivation

Dataflow Systems




Graph Systems



GraphX Motivation


- Difficult to Program and Use
- Users must *Learn, Deploy, and Manage* multiple systems



- Leads to brittle and often complex interfaces

And Inefficient


- Extensive **data movement** and duplication across the network and file system



- Limited reuse of internal data-structures across stages

The GraphX Unified Approach

New API Blurs the distinction between
Tables and Graphs

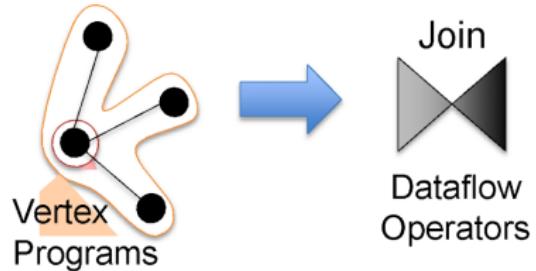
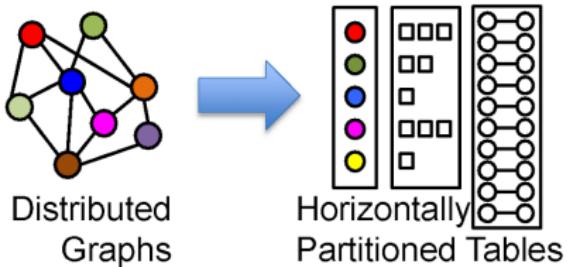


New System Combines Data-Parallel
Graph-Parallel Systems



Enables users to easily and efficiently express the entire graph analytics pipeline

Representation



- Plus optimizations:
 - Distributed join optimization
 - Materialized view maintenance

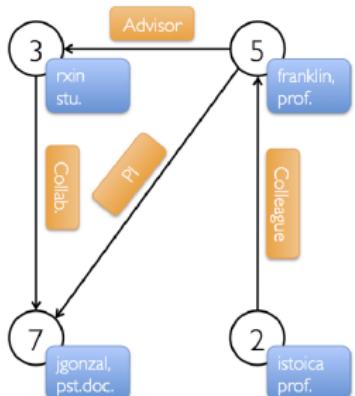
Graph modeling in GraphX

- The property graph is parameterized over the vertex (VD) and edge (ED) types

```
class Graph[VD, ED] {  
    val vertices: VertexRDD[VD]  
    val edges: EdgeRDD[ED]  
}
```

- Graph[(String, String), String]

Property Graph Vertex Table



Id	Property (V)
3	(rxin, student)
7	(jgonzal, postdoc)
5	(franklin, professor)
2	(istoica, professor)

Edge Table

SrcId	DstId	Property (E)
3	7	Collaborator
5	3	Advisor
2	5	Colleague
5	7	PI

Creating a Graph (Scala)


```
type VertexId = Long

val vertices: RDD[(VertexId, String)] =
  sc.parallelize(List(
    (1L, "Alice"),
    (2L, "Bob"),
    (3L, "Charlie")))

class Edge[ED]{
  val srcId: VertexId,
  val dstId: VertexId,
  val attr: ED)

  val edges: RDD[Edge[String]] =
    sc.parallelize(List(
      Edge(1L, 2L, "coworker"),
      Edge(2L, 3L, "friend")))

  val graph = Graph(vertices, edges)
```



Hello world in GraphX

```
$ spark*/bin/spark-shell
scala> val inputFile = sc.textFile("hdfs:///tmp/graph/1.txt")
scala> val edges = inputFile.flatMap(s \$\implies$ {
val l = s.split("\t");
l.drop(1).map(x \$\implies$ (l.head.toLong, x.toLong))
})
scala> val graph = Graph.fromEdgeTuples(edges, "")
scala> val result = graph.collectNeighborIds(EdgeDirection.Out).ma
println("Hello world from the: " + x._1 + " : " + x._2.mkString(" ")
scala> result.collect() // don't try this @home
```

```
Hello world from the: 1 :
Hello world from the: 2 : 1 3
Hello world from the: 3 : 1 2
```

Spark Table Operators

- GraphX **Table** (RDD) operators are inherited from Spark:
- map
- filter
- groupBy
- sort
- union
- join
- leftOuterJoin
- rightOuterJoin
- reduce
- count
- fold
- reduceByKey
- groupByKey
- cogroup
- cross
- zip
- sample
- take
- first
- partitionBy
- mapWith
- pipe
- save
- ...


Graph Operators (Scala)

```
class Graph [ V, E ] {  
    def Graph(vertices: Table[ (Id, V) ],  
              edges: Table[ (Id, Id, E) ])  
    // Table Views -----  
    def vertices: Table[ (Id, V) ]  
    def edges: Table[ (Id, Id, E) ]  
    def triplets: Table [ ((Id, V), (Id, V), E) ]  
    // Transformations -----  
    def reverse: Graph[V, E]  
    def subgraph(pV: (Id, V) $\implies$ Boolean,  
                pE: Edge[V,E] $\implies$ Boolean): Graph[V,E]  
    def mapV(m: (Id, V) $\implies$ T ): Graph[T,E]  
    def mapE(m: Edge[V,E] $\implies$ T ): Graph[V,T]  
    // Joins -----  
    def joinV(tbl: Table [(Id, T)]): Graph[(V, T), E ]  
    def joinE(tbl: Table [(Id, Id, T)]): Graph[V, (E, T)]  
    // Computation -----  
    def mrTriplets(mapF: (Edge[V,E]) $\implies$ List[(Id, T)],  
                  reduceF: (T, T) $\implies$ T): Graph[T, E]  
}
```

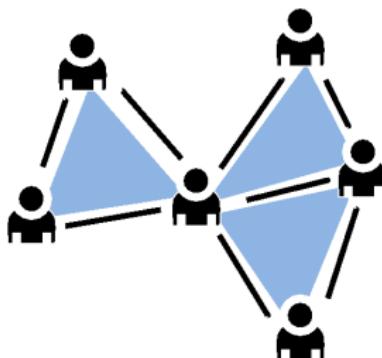
Built-in Algorithms (Scala)

```
def pageRank(tol: Double): Graph[Double, Double]
def triangleCount(): Graph[Int, ED]
def connectedComponents(): Graph[VertexId, ED]
def stronglyConnectedComponents(numIter:Int):Graph[VertexID,ED]
```

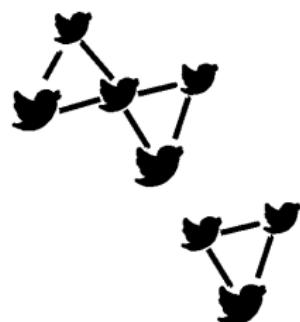
PageRank



Triangle Count



Connected Components



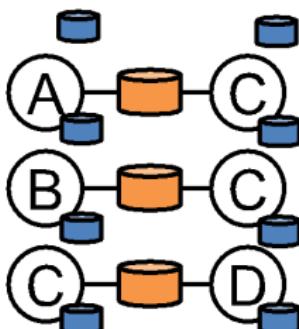
Triplets Join Vertices and Edges

- Triplets capture Gather-Scatter pattern from specialized graph processing systems (like Giraph)
- **Triplets** operator joins vertices and edges

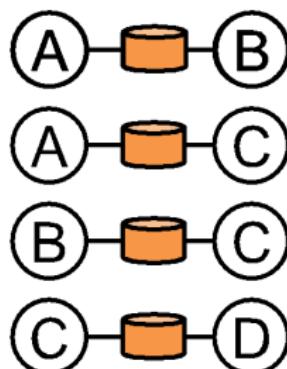
Vertices



Triplets

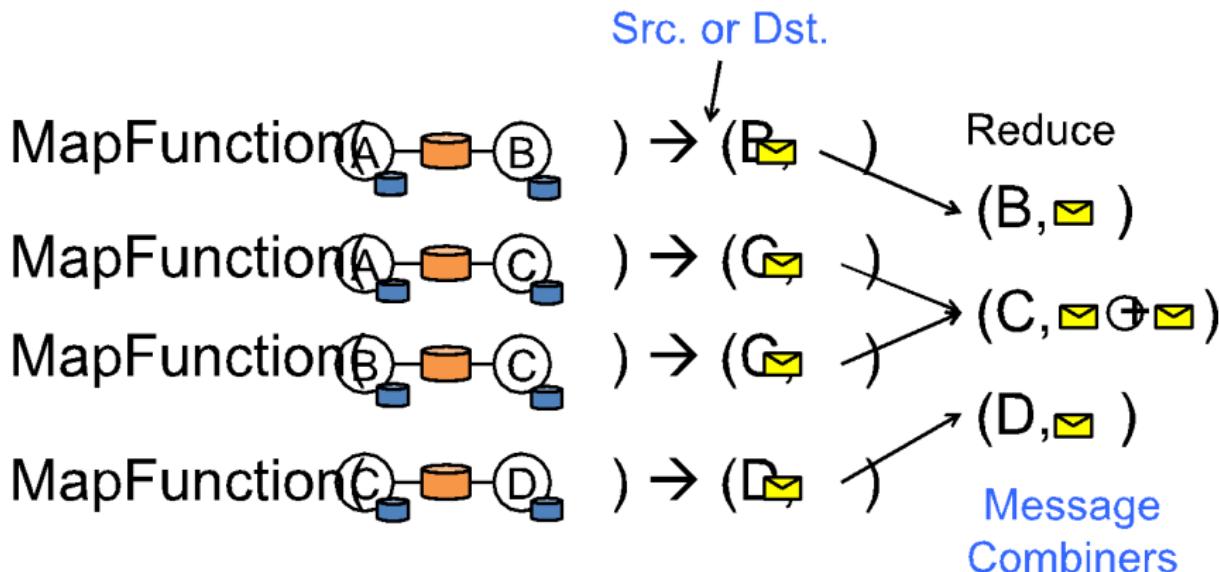


Edges




MapReduce Triplets

Map-Reduce triplets collect information about the neighborhood of each vertex:



Performance Comparisons

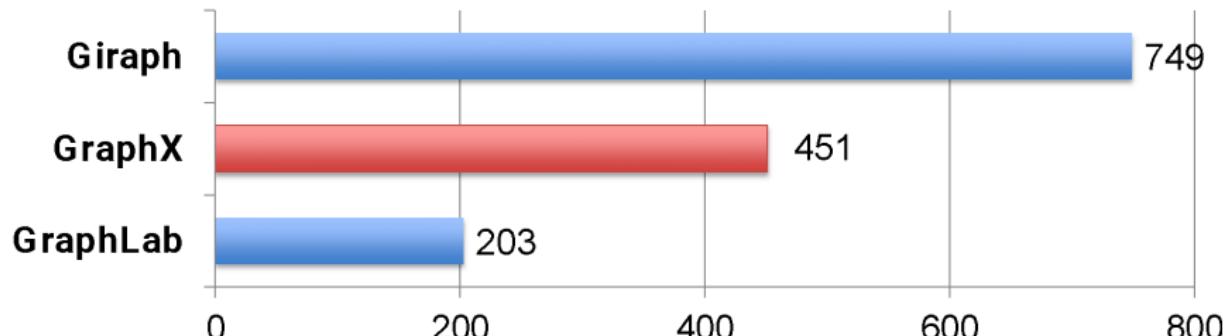
Live-Journal: 69 Million Edges



GraphX is roughly **3x slower** than GraphLab

GraphX scales to larger graphs

Twitter Graph: 1.5 Billion Edges

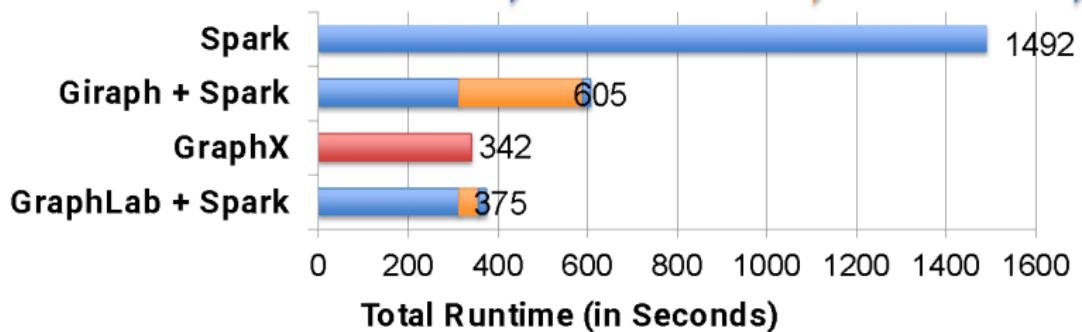



- GraphX is roughly 2x slower than GraphLab - Scala + Java overhead: Lambdas, GC time, ... - No shared memory parallelism: 2x increase in communication

But, a Small Pipeline in GraphX

Timed end-to-end GraphX is faster than GraphLab (and Giraph)

Raw Wikipedia Hyperlinks PageRank Top 20 Pages



Giraph vs. GraphX

- **Giraph**

- An unconstrained BSP framework
- Specialized fully mutable, dynamically balanced in-memory graph representation
- Procedural, vertex-centric programming model
- Part of Hadoop ecosystem

- **GraphX**

- An RDD framework
- Graphs are “views” on RDDs and thus immutable
- Functional-like, “declarative” programming model
- Genuine part of Spark ecosystem