



UMD DATA605 - Big Data Systems

12.3: Graph Data Processing

- **Instructor:** Dr. GP Saggese, gsaggese@umd.edu

- *Typical graph analysis tasks*
- Graph Data Processing Systems

Queries vs Analysis Tasks

- **Queries**

- Explore data
- Result is small graph portion (often a node)
- Challenges
 - Minimize explored graph portion
 - Use indexes (auxiliary data structures)

- **Analysis tasks**

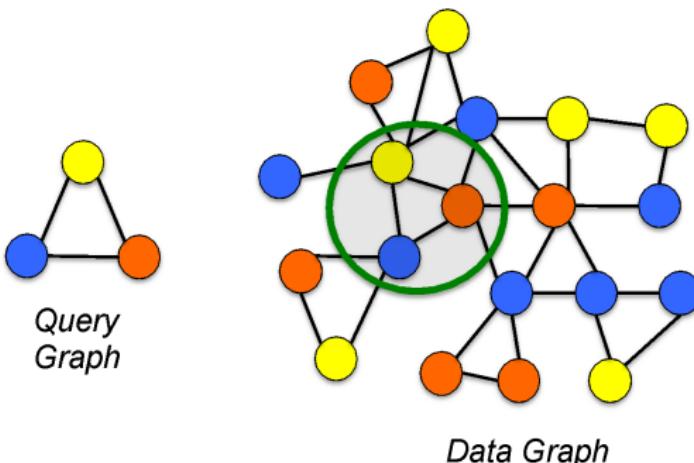
- Process entire graph
- Challenges
 - Handle large data efficiently
 - Parallelize if data doesn't fit in memory/disk

Graph Algorithms

- **Graph algorithms** are general algorithms applicable to various graph types
 - Network flows (e.g., max flow, min cut)
 - Spanning trees (e.g., minimal connectivity)
 - ...
- **Applications**
 - Logistics and supply chain optimization
 - Electric grid design
 - Telecommunications (e.g., bandwidth allocation)

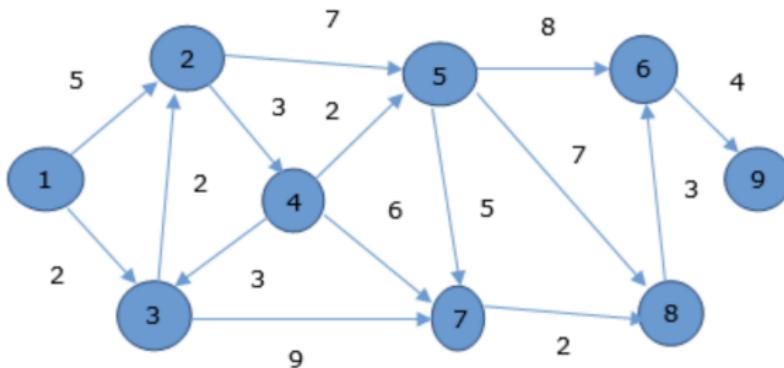
Subgraph Matching

- **Subgraph Matching** = find matching instances of a small pattern in a larger graph
 - Patterns are typically small and fixed
 - Matching can be *approximate*
- **Applications**
 - Fraud detection (e.g., detecting fraudulent transaction patterns)
 - Bioinformatics (e.g., finding protein interaction motifs)
 - Social network analysis (e.g., community pattern detection)



Shortest Path Queries

- **Shortest Path Queries** = find the shortest or optimal path between two nodes
 - May consider edge weights (e.g., distance, cost)
- **Applications**
 - GPS navigation (e.g., Google Maps)
 - Network routing (e.g., Internet traffic routing)
 - Robotics (e.g., path planning)

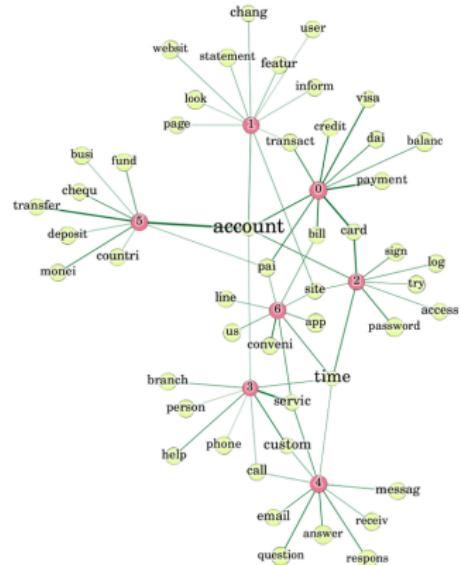


Reachability

- **Reachability** = determine if a path exists between two nodes
 - May include constraints (e.g., edge types, direction)
- **Applications**
 - Access control (e.g., "can user X reach resource Y?")
 - Dependency analysis (e.g., package installation)
 - Workflow engines (e.g., task progression)

Keyword Search

- **Keyword Search** = find smallest subgraph containing all specified keywords
- **Applications**
 - Knowledge graphs (e.g., question answering)
 - Enterprise search (e.g., documents linked by shared entities)
 - Academic citation networks



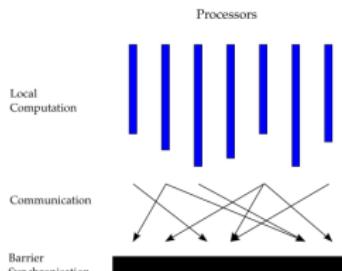
Historical Queries

- **Historical Queries** = find nodes with similar evolution or temporal behavior
 - Typically on dynamic or time-stamped graphs
- **Applications**
 - Stock market analysis (e.g., similar price movement)
 - Social media trends (e.g., user behavior over time)
 - Epidemiology (e.g., disease spread patterns)

- Typical graph analysis tasks
- ***Graph Data Processing Systems***

Bulk Synchronous Parallel Model

- BSP is a model for parallel computation using synchronized steps
- Computation is organized into supersteps
 - Local computation
 - Message passing
 - Parallel processors to exchange data
 - Messages are sent during a superstep but delivered in the next
 - Synchronization barriers
 - Ensure all processors complete a superstep before moving on
- Pros
 - Deterministic and predictable behavior
 - Simple abstraction for reasoning about parallelism
 - Suitable for graph computations with regular communication patterns
- Cons
 - Global synchronization can cause idle time (stragglers)
 - Inefficient for highly irregular workloads or dynamic computation patterns



Pregel

- **Pregel** is a graph processing framework inspired by the BSP model
 - Pregel = Parallel + Graph + Google
 - Designed for large-scale distributed graph computations
 - Uses a vertex-centric programming model
 - Each vertex performs computation independently in supersteps
 - Motto: “*Think like a vertex*”
 - Vertices
 - Send messages to neighbors
 - Update state based on received messages
 - Can vote to halt; computation ends when all are inactive
- **PageRank example**
 - Each vertex divides its rank among neighbors
 - In each superstep, vertices update rank from incoming messages
- **Connected components example**
 - Each vertex propagates smallest ID seen to neighbors
 - Repeats until no changes occur

Apache Giraph

- **Apache Giraph** is an open-source implementation of Pregel
 - Written in Java for high scalability
- Suitable for batch-processing large graphs
- **Integrates with Hadoop ecosystem**
 - Runs on Hadoop MapReduce
 - Uses HDFS for storage and YARN for resource management
- **Example use cases:**
 - Social network analysis (e.g., friend recommendation)
 - Web graph analysis (e.g., PageRank)
 - Biological network exploration (e.g., gene interaction)



Apache Spark GraphX



- **GraphX** is Spark's API for graph-parallel computation
- Seamless integration with Spark ecosystem and pipelines
- Graphs represented using RDDs
 - Vertices and edges as distributed collections
- Includes a Pregel API
 - Enables iterative message-passing computations
 - Preserves immutability and functional programming style
- **Strengths vs traditional graph DBs**
 - Scalable to very large graphs
 - Efficient for batch analytics and machine learning workflows
- **Limitations**
 - Less suitable for real-time queries and transactional workloads
 - Requires understanding of Spark internals for tuning performance