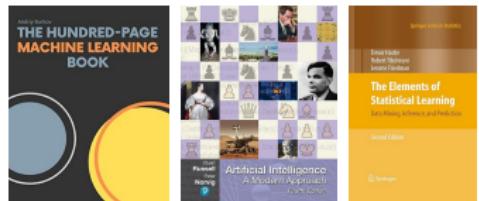


Lesson 02.2: Machine Learning Paradigms

Instructor: Dr. GP Saggese - gsaggese@umd.edu

References:

- Burkov: “*The Hundred-Page Machine Learning Book*” (2019)
- Russell et al.: “*Artificial Intelligence: A Modern Approach*” (4th ed, 2020)
- Hastie et al.: “*The Elements of Statistical Learning*” (2nd ed, 2009)



- ***Machine Learning Paradigms***

- Major Paradigms
- Some Examples
- Machine Learning in Practice
- Pipeline Organization

- Machine Learning Paradigms
 - *Major Paradigms*
 - Some Examples
 - Machine Learning in Practice
 - Pipeline Organization

Machine Learning Paradigms with Examples (1/3)

- **Supervised Learning**

- Learn from labeled data to predict labels for new inputs
- E.g., image classification using ResNet on ImageNet

- **Unsupervised Learning**

- Discover hidden patterns or structure in unlabeled data
- E.g., K-means clustering for customer segmentation

- **Reinforcement Learning**

- Learn through interaction with an environment, receiving rewards/punishments
- E.g., deep Q-Learning for playing Atari games

- **Self-Supervised Learning**

- Generate pseudo-labels from unlabeled data to pre-train models
- E.g., BERT (Masked Language Modeling)

- **Semi-Supervised Learning**

- Combine small labeled data with large unlabeled data to improve performance
- E.g., named entity recognition (NER) using annotated sentences with entity tags combined with many raw text documents



Machine Learning Paradigms with Examples (2/3)

- **Online Learning**

- Learn incrementally from a stream of data in real time
- E.g., online logistic regression for click-through rate prediction

- **Multi-Task Learning**

- Train simultaneously a model to perform multiple related tasks
- E.g., learn sentiment analysis and question answering

- **Meta-Learning**

- “Learning to learn”: adapt quickly to new tasks using prior experience
- E.g., a model can be fine-tuned quickly on a new task using just a few gradient steps

- **Zero-Shot / Few-Shot Learning**

- Generalize to new tasks with no or few labeled examples
- E.g., GPT-4 solving tasks with zero-shot prompting

- **Active Learning**

- The model selects the most informative samples to be labeled by an oracle (e.g., a human)
- E.g., pick samples where the model is least confident to get more examples



Machine Learning Paradigms with Examples (3/3)

- **Federated Learning**

- Train models across decentralized devices without sharing raw data
- E.g., fraud detection or credit scoring across banks

- **Evolutionary Learning**

- Optimize model structures or parameters using evolutionary algorithms inspired by natural selection and genetics
- Gradient free, global search, variable length inputs
- E.g., genetic algorithms

- **Curriculum Learning**

- Train models on easier tasks first, gradually increasing difficulty
- E.g., training in robotic control simulations

- **Multi-Agent Learning**

- Multiple agents learn and interact in shared environments, often in game-theoretic settings (e.g., competition, collaboration)
- E.g., AlphaStar to play StarCraft II

- Machine Learning Paradigms
 - Major Paradigms
 - ***Some Examples***
 - Machine Learning in Practice
 - Pipeline Organization

Supervised Learning

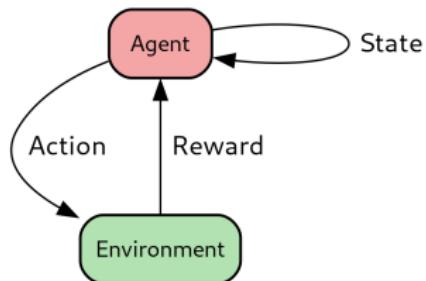
- Learn a function $f : X \rightarrow Y$ that maps inputs to correct outputs
 - Training examples (\underline{x}, y) with pairs of inputs and correct outputs
 - Requires labeled data for training
 - Measure performance with error on a separate test set
- **Classification:** output is a discrete label, e.g.,
 - Spam vs Not Spam
 - Digit recognition 0, 1, ...
 - Sentiment analysis Pos, Neg, Neutral
- **Regression:** output is a continuous value, e.g.,
 - House prices given features like size and location
 - House demand
 - Stock prices
- **Common algorithms:**
 - Linear Regression
 - Decision Trees
 - K-nearest neighbors
 - Neural Networks
 - ...

Unsupervised Learning

- Learn from data **without** labeled outputs
 - Goal: discover patterns or structure in the data
 - No explicit feedback signal
 - Evaluation can be qualitative
- **Main techniques:**
 - **Clustering:** Group similar examples, e.g.,
 - Customer segmentation
 - Grouping news articles by topic without knowing the topics
 - **Dimensionality Reduction:** Reduce number of variables with PCA while preserving structure
 - E.g., visualize high-dimensional data in 2D
 - **Density Estimation:** Estimate probability distribution of data
 - E.g., anomaly detection in server logs
 - **Association Rule Learning:** Discover interesting relations between variables
 - E.g., market basket analysis (e.g., “people who buy X also buy Y”)
- **Common algorithms:**
 - K-means
 - PCA
 - Autoencoders

Reinforcement Learning

- Learn by **interacting with an environment** to **maximize cumulative reward**
 - Learn policy $\pi(s) \rightarrow a$ that maximizes expected reward
 - Trade-off between exploration (trying new actions) and exploitation (using known good actions)
 - Environments provide clear rules and feedback (win/loss/reward)
 - Often involve physical simulation or real-world interaction
- **Core elements:**
 - Agent: Learner and decision maker
 - Environment: Everything the agent interacts with
 - State s
 - Action a
 - Reward r
- **Algorithms:**
 - Q-learning
 - Policy Gradient methods



Reinforcement Learning: Examples

- In **game playing**, learn strategies through trial and error
 - E.g., AlphaGo mastering the game of Go
- In **robotics**, learn control policies for movement and manipulation
- In **autonomous driving**, learn safe and efficient driving behaviors
- In **resource management**, optimize allocation of limited resources over time
 - E.g., data center cooling or CPU job scheduling
- In **personalized recommendations**, adapt suggestions based on user interaction
 - E.g., newsfeed ranking adjusting based on user clicks
- In **healthcare**, optimize treatment plans over time

- Machine Learning Paradigms
 - Major Paradigms
 - Some Examples
 - ***Machine Learning in Practice***
 - Pipeline Organization

Machine Learning Flow (1/2)

- **Question**
 - E.g., "How can we predict house prices?"
- **Input data**
 - E.g., historical data of house sales
- **Features**
 - E.g., number of bedrooms, location, square footage
- **Algorithm**
 - E.g., linear regression, decision trees
- **Parameters**
 - E.g., learning rate, number of trees in a random forest
- **Evaluation**
 - E.g., accuracy, precision, recall

Machine Learning Flow (2/2)



- Not all phases are equally important!
 - Question > Data > Features > Algorithm
- Clarity of the **question** impacts project success
- Quality and relevance of **data** are crucial for performance
- Proper **feature** selection simplifies the model and improves accuracy
- **Algorithm** is often less important (contrary to popular belief!)

Question

- **Make the question concrete and precise**
 - Define the problem clearly
 - Specify inputs and expected outputs
 - Align question with business or research objectives
 - E.g.,:
 - **Bad:** “How can we improve sales?”
 - **Good:** “What factors most significantly impact sales of product X in region Y during season Z?”
- Formulating question is **the most important part** of the machine learning problem
 - Misunderstanding leads to:
 - Solving the wrong problem
 - Collecting wrong data
 - ...
- “*If I were given one hour to save the planet, I would spend 59 minutes defining the problem and one minute resolving it*” (Einstein)

Input Data

- Ensure **data is specific to prediction** goal
 - E.g., use known movie ratings to predict unseen movie ratings from the same population
 - Training set \approx test set
- Relationship between data and prediction goal is **not always direct**
 - E.g., interested in prices but predict supply and demand instead
- Poor-quality data leads to inaccurate predictions
 - “*Garbage in - garbage out*”
- Recognize **when data is insufficient** for valid answers
 - *“Combination of data and desire for an answer does not ensure a reasonable answer can be extracted”* (John Tukey)
- **More data vs better models**
 - Meta-studies show difference between generic and best model is like 5%
 - *“It's not who has the best algorithm that wins. It's who has the most data”* (Google researcher)
 - *“Every time I fire a linguist, the performance of the speech recognizer goes up”* (IBM researcher in speech recognition)

Features

- **Features** provide high-level information about inputs
 - E.g., use intensity and symmetry for scanned numbers instead of raw bit maps
- **Characteristics of good features:**
 1. Perform data compression
 2. Retain relevant information
 3. Often created with expert knowledge
- **Common mistakes in feature building:**
 1. Automating feature selection may lead to overfitting
 - Black box predictions can be accurate but stop working anytime
 - E.g., Google Flu's unclear feature-model link
 2. Ignoring data-specific quirks
 - E.g., mislabeling outliers
 3. Unnecessarily discarding information

Models

- Best models are:
 - **Interpretable**
 - Allow users to understand and trust the model's decisions
 - E.g., decision trees are appropriate in medical studies since they produce a "reasoning"
 - **Simple**
 - Easier to implement and maintain
 - Reduces the risk of overfitting
 - **Accurate**
 - Often accuracy is traded off for remaining characteristics
 - E.g., accuracy vs interpretability, accuracy vs speed
 - **Fast**
 - Reduce computational costs to train and test
 - E.g., real-time applications
 - **Scalable**
 - Can handle large datasets efficiently (growing data and user bases)
 - E.g., in the Netflix prize, Netflix didn't end up implementing the best algorithm since it wasn't scalable enough

- Machine Learning Paradigms
 - Major Paradigms
 - Some Examples
 - Machine Learning in Practice
 - *Pipeline Organization*

How Are Machine Learning Systems Organized?

- Machine learning systems organized in a **pipeline**
 1. **Break down** problem into sub-problems
 2. **Solve** problems one at a time
 3. **Combine** solutions to solve initial problem
- Performance p of a pipeline:

$$p_{\text{system}} = \sum_i p_i \cdot \alpha_i$$

where:

- p_i : Performance of each stage
- α_i : Importance of each stage

Example of Photo OCR System

- **Goal:** build systems to read text in a picture
 - OCR = “Optical Character Recognition”
- **Stages of ML pipeline for OCR:**
 - Text detection: find text areas in the picture
 - Character segmentation: split text into letter boxes
 - E.g., h e l l o
 - Character classification: classify characters individually
 - Spelling correction: fix text errors using context
 - E.g., hell0 corrected to hello
- **Sliding window approach**
 - **Problem:** Unknown text location and size
 - Text detection
 - Train classifier to recognize letters vs non-letters
 - Scan image in two directions with different sizes for text
 - Evaluate classifier cheaper than training
 - Create text likelihood map (e.g., heatmap) using classifier probabilities
 - Enclose text areas in boxes
 - Discard boxes not fitting aspect ratio (valid text width > height)
 - Character segmentation
 - Use sliding window classifiers to find “breaks” between characters

Getting More Data

- The ideal recipe for ML is: “low-bias algorithm + train on massive amount of data”
- Use learning curves to make sure we are taking advantage of more data
- The **problem** is getting large amount of data
- Always ask yourself: *“how much work is to get 10x more data than we currently have?”*
- Often it is not that difficult:
 1. Artificial data
 - E.g., synthesize or amplify data set
 2. Collect and label by hand
 - E.g., crowd sourcing like Amazon Mechanical Turk

OCR Pipeline: Example of Artificial Data Synthesis

- How can we increase data set size?
 1. Synthesize data set
 - Use font libraries to generate large training sets
 - Paste characters against random backgrounds
 - Apply scaling, distortion, adding noise, etc
 2. Amplify a data set
 - Add examples by warping/distorting existing examples
- Transformations and noise should be specific to the application domain
 - E.g., Gaussian noise is not always appropriate

Ceiling Analysis for ML Pipeline

- The most valuable resource is time
 - You work on optimization for months
 - Find out that the optimization doesn't make much difference
- **Problem:** On which part of the pipeline should time/resource be spent?
- **Solution:** Ceiling analysis to Analyze pipeline performance
 - Single number for system performance
 - E.g., accuracy for OCR
 - For each component:
 - Mock component with "oracle" (always give correct output)
 - Leave others untouched
 - Compute pipeline performance
 - Identify critical components by estimating performance improvement with 10% component improvement
 - Measure, don't guess!