


Lesson 02.4: ML Techniques - Model Learning

Instructor: Dr. GP Saggese - gsaggese@umd.edu

References:


- Burkov: “*The Hundred-Page Machine Learning Book*” (2019)
- Russell et al.: “*Artificial Intelligence: A Modern Approach*” (4th ed, 2020)
- Hastie et al.: “*The Elements of Statistical Learning*” (2nd ed, 2009)



- *Learning Algorithms*


The Problem of Minimizing a Function

- **Goal:** minimize scalar function $J(\underline{w})$ of n -variables \underline{w}
 - E.g., in-sample error $E_{in}(\underline{w})$
- **Solutions:**
 1. Analytical solution
 - Impose the gradient of $J(\underline{w})$ to equal 0
 - Find a closed-form solution for \underline{w}^*
 2. Numerical solution:
 - Use an iterative method to update \underline{w} to reach the minimum value of $J(\underline{w})$
 - E.g., gradient descent
 - It works even if there is an analytical solution



Gradient Descent: Intuition

- **Problem:**
 - You are on a hilly surface and we want to walk down to the bottom of the hill
- **Solution:**
 - At each point:
 - You look around
 - You move a step in the direction where the surface is steepest
 - You keep doing until we reach the bottom
- **Gradient descent**
 - Is a general technique for minimizing twice-differentiable functions
 - In general, converge to a local minimum
 - If $J(\mathbf{w})$ is convex, converge to the global minimum
 - E.g., logistic regression and linear models



Gradient descent with fixed learning rate (1/3)

- Consider the contour plot of a function
- **Start** from a point $\underline{w}(0)$ (random, the origin, ...)
- **At each step**
 - Move a fixed amount η in weight space (fixed learning rate):

$$\underline{w}(t + 1) = \underline{w}(t) + \eta \hat{\underline{v}}$$

where $\hat{\underline{v}}$ is a unit vector

- Pick $\hat{\underline{v}}$ to move to a value of $E_{in}(\underline{w})$ as negative as possible
- The **change for E_{in}** is:

$$\begin{aligned}\Delta E_{in} &= E_{in}(\underline{w}(t + 1)) - E_{in}(\underline{w}(t)) \\ &= E_{in}(\underline{w}(t) + \eta \hat{\underline{v}}) - E_{in}(\underline{w}(t)) \quad (\text{replacing the expression of } \underline{w}(t + 1)) \\ &\approx \eta \nabla E_{in}(\underline{w}(t))^T \hat{\underline{v}} + O(\eta^2) \quad (\text{using Taylor expansion})\end{aligned}$$

- Gradient descent keeps only $O(\eta)$ the term and ignores the rest
- Conjugate gradient considers up to $O(\eta^2)$ and ignores higher infinitesimals



Gradient Descent with Fixed Learning Rate (2/3)

- You have:

$$\Delta E_{in} \approx \eta \nabla E_{in}(\underline{\mathbf{w}}(t))^T \underline{\hat{\mathbf{v}}}$$

- The **minimal value of the scalar product** ΔE_{in} :

- Happens when $\underline{\hat{\mathbf{v}}} = -\frac{\nabla E_{in}(\underline{\mathbf{w}}(t))}{\|\nabla E_{in}(\underline{\mathbf{w}}(t))\|}$
- Is $-\eta \|\nabla E_{in}(\underline{\mathbf{w}}(t))\|$

- The **change in weights** $\Delta \underline{\mathbf{w}}$ is:

$$\begin{aligned}\Delta \underline{\mathbf{w}} &= \underline{\mathbf{w}}(t+1) - \underline{\mathbf{w}}(t) \\ &= \eta \underline{\hat{\mathbf{v}}} \\ &= -\eta \frac{\nabla E_{in}(\underline{\mathbf{w}}(t))}{\|\nabla E_{in}(\underline{\mathbf{w}}(t))\|}\end{aligned}$$

- Called "**gradient descent**" since it descends along the gradient of the function to optimize

Gradient Descent with Fixed Learning Rate (3/3)

- Each j component of the weight $\underline{\mathbf{w}}$ is updated with the partial derivative with respect to that coordinate:

$$\underline{\mathbf{w}}(t+1) = \underline{\mathbf{w}}(t) - \eta \frac{\nabla E_{in}(\underline{\mathbf{w}}(t))}{\|\nabla E_{in}(\underline{\mathbf{w}}(t))\|}$$
$$w_j(t+1) = w_j(t) - \eta \frac{1}{\|\nabla E_{in}(\underline{\mathbf{w}}(t))\|} \frac{\partial E_{in}(\underline{\mathbf{w}})}{\partial w_j}$$




- The update of all components should be **simultaneous**, i.e., computed at once
- A step of the optimization when we update the solution (weights) is called **epoch**

Gradient Descent: Stopping Criteria

- **In theory**, stop when $\Delta E_{in} = \underline{0}$
 - Numerically, this might not occur
- **In practice**, stop when:
 - Variation of E_{in} is smaller than threshold $\Delta E_{in} < \theta$
 - Reached a certain number of iterations
- **Monitoring gradient descent**
 - In theory, only compute derivatives of function $J(\underline{w})$ to optimize
 - In practice, monitor progress by recomputing cost function $J(\underline{w})$ periodically to ensure it decreases

Gradient Descent with Fixed Learning Rate

- Consider a 1D convex function
 - If η is too small:
 - Linear approximation of E_{in} is effective
 - Many steps needed to converge to minimum
 - If η is too large:
 - Linear approximation fails (higher terms affect values)
 - It “bounces around”



- Idea:** vary learning rate η during gradient descent
 - Smaller learning rates may find a better minimum
 - Reduce η as a function of iterations

Cons: introduces an additional parameter to tune

Gradient Descent with Variable Learning Rate

- In **gradient descent with fixed learning rate** (constant change in weight space), use:

$$\Delta \underline{w} = -\eta \frac{\nabla J}{\|\nabla J\|}$$

- **To converge quickly:**

- Move fast (large change) in weight space (large η) when surface is steep (large gradient)
- Move slow (small change) in weight space (small η) near minimum to avoid bouncing (small gradient)
- Ideally, η should increase with slope: $\eta \propto \|\nabla J\|$

- **Gradient descent with variable learning rate:**

$$\Delta \underline{w} = -\eta \nabla J$$

Feature Scaling in Gradient Descent

- **Problem:** Different gradient components have different errors due to numerical approximation, causing bouncing
 - Unscaled features lead to slow, unstable convergence
 - E.g., feature ranging from 1 to 1000 vs. 0.01 to 1 causes inefficient updates
- **Solution:** Converge faster by scaling features to the same range
 - Min-max scaling
 - Standardization (transforms feature values to mean 0, standard deviation 1)

Issues with Batch Gradient Descent

- Consider squared error with n samples:

$$E_{in}(\underline{\mathbf{w}}) \triangleq \frac{1}{n} \sum_{i=1}^n (h_{\underline{\mathbf{w}}}(\underline{\mathbf{x}}_i) - y_i)^2$$

- Batch Gradient Descent (BGD)** updates each weight component:

$$\underline{\mathbf{w}}(t+1) = \underline{\mathbf{w}}(t) - \eta \frac{\nabla E_{in}}{\|\nabla E_{in}\|}$$

- Coordinate update for squared error:

$$w_j(t+1) = w_j(t) - \eta \frac{2}{n} \sum_{i=1}^n (h_{\underline{\mathbf{w}}}(\underline{\mathbf{x}}_i) - y_i) \frac{\partial h_{\underline{\mathbf{w}}}(\underline{\mathbf{x}}_i)}{\partial w_j}$$

- Cons:** gradient descent with many training examples (e.g., $N = 10^6$)
 - Is computationally expensive, requiring gradient evaluation from all examples for one update
 - Requires storing all data in memory

Stochastic Gradient Descent

- Idea of **Stochastic Gradient Descent (SGD)**: update the weights only for one training example picked at random
- Algorithm**

- Pick one (\underline{x}_n, y_n) at a time from the available examples
- Compute $\nabla e(h(\underline{x}_n), y_n)$ to update the weights: $\Delta \underline{w} = -\eta \nabla e$
- Update the weight considering only one random example:

$$w_j(t+1) = w_j(t) - \eta \frac{2}{n} (h_{\underline{w}}(\underline{x}_t) - y_t) \frac{\partial h_{\underline{w}}(\underline{x}_t)}{\partial w_j}$$

- Cons:** in Stochastic Gradient Descent

- Path in weight space is random
- Oscillates around local minimum

- Why SGD works**

- ∇e is a function of a random var \underline{x}_n
- The average direction of SGD is the same direction as batch version


$$\mathbb{E}[\nabla e] = \frac{1}{N} \sum \nabla e(h(\underline{x}_n), y_n) = \nabla \frac{1}{N} \sum e(h(\underline{x}_n), y_n) = \nabla E_{in}$$

Mini-Batch Gradient Descent

- **Mini-Batch Gradient Descent** brings together characteristics of both Batch and Stochastic Gradient Descent
 - Use b examples to make an update to the current weight
 - b represents the batch size
 - A common choice is $b = 32$ or $b = 64$
- Mini-batch GD offers a balance between SGD noisiness and full-batch approaches, using small, random data samples for updates

On-Line Learning and Gradient Descent

- In many applications, **continuous stream** of training examples → requires updating the model
 - In real-time systems, adapt model with new data points without full retraining
 - Handle variation in underlying process dynamics
- Stochastic GD and mini-batch GD **suitable for online learning**
 - Update model one example at a time
 - Discard examples for “compressed” model representation
 - Useful for large data streams where storing every data point is impractical
 - E.g., in stock market prediction
 - E.g., in training a language model on live chat data, discard older conversations after updates




SGD vs BGD vs Mini-Batch

- To update the weights:
 - **BGD (batch gradient descent)** uses all the training examples
 - **SGD (stochastic gradient descent)** uses a single random training example
 - **Mini-batch GD** uses a random subset of training examples

Aspect	Batch Gradient Descent	Stochastic Gradient Descent
<i>Computation</i>	Uses all examples	One example at a time
<i>Memory</i>	Requires all examples in memory	Require less memory
<i>Randomization</i>	More likely to terminate in flat regions	Avoid local minima due to randomness
<i>Regularization</i>	No implicit regularization	Oscillations act as regularization
<i>Parallelization</i>	Can be parallelized	Less parallel-friendly
<i>Online Learning</i>	Not suitable	Suitable for online learning

Map-Reduce for Batch Gradient Descent

- Batch GD (and many learning algorithms) can be expressed in map-reduce
- **In map-reduce**
 - **Map step:** use k machines to parallelize summation
 - **Reduce step:** send k partial sums to a single node to accumulate result



Coordinate Descend

- **Idea:** minimize $J(x_0, \dots, x_n)$ by optimizing along one direction x_i at a time
 - Instead of computing all derivatives
- **Algorithm**
 - Pick a random starting point $\underline{w}(0)$
 - Pick a random order for the coordinates $\{x_i\}$
 - Find the minimum along the current coordinate (1D optimization problem)
 - Move to the next coordinate x_{i+1}
 - The sequence of $\underline{w}(t)$ is decreasing
 - A minimum is found if there is no improvement after one cycle of scanning all coordinates
 - The minimum is local

Gradient Descent vs Pseudo-Inverse

- For linear models, you can use pseudo-inverse or gradient descent to find optimal $\underline{\mathbf{w}}^*$
- **Gradient descent**
 - Choose learning rate η
 - Requires many iterations
 - Monitor stopping criteria, oscillations
 - Effective for many features P
- **Pseudo-inverse**
 - No parameter selection
 - Converges in one iteration
 - Computes a $P \times P$ matrix $(\underline{\mathbf{X}}^T \underline{\mathbf{X}})^{-1}$
 - Complexity of matrix inversion $O(P^3)$
 - For $P \approx 10,000$, prefer gradient descent