



MSML610: Advanced Machine Learning

Lesson 09.3: Multi-armed Bandits

Instructor: Dr. GP Saggese, gsaggese@umd.edu

References:

- *Introduction*
- Algorithms
- Bayesian Approaches

What are Multi-Armed Bandits?

- Imagine you're in a casino facing K slot machines (aka "one-armed bandits")
 - Each machine has a different unknown probability of winning
 - Limited budget of plays
 - *Goal:* Maximize total winnings
- **Multi-Armed Bandit (MAB)** is a sequential decision-making problem where:
 - At each time step $t = 1, 2, \dots, T$, choose one action (arm) from K options
 - After choosing arm i , receive a random reward $r_t \sim P_i$
 - Reward distributions P_1, \dots, P_K are unknown
 - Learn which arms are best while collecting rewards
- Sequential decisions with uncertain outcomes and limited feedback
- **Key Challenge:** Balance:
 - *Exploration:* Try different arms to learn their reward distributions
 - *Exploitation:* Play the arm you believe is best

Why Study Bandits?

- **Fundamental Problem in Sequential Decision-Making**
 - Core model for learning under uncertainty
 - Simplest non-trivial reinforcement learning setting
 - Rich theoretical framework with practical algorithms
- **Tractable Analysis**
 - Stateless: No long-term consequences of actions
 - Clean mathematical formulation
 - Provable regret bounds and optimality results
 - Serves as building block for more complex RL problems
- **Wide Applicability**
 - Online advertising (which ad to show?)
 - Recommendation systems (which movie to suggest?)
 - Clinical trials (which treatment to test?)
 - Website optimization (A/B testing)
 - Resource allocation (which server to query?)

Applications: Real-World Decision-Making

- **Online Advertising**

- Platform must choose which ad to display
- Each ad click generates revenue
- Must learn which ads perform best
- Trade-off: Show proven ads vs. test new ones

- **Clinical Trials**

- Allocate patients to treatment arms
- Each treatment has unknown efficacy
- *Ethical imperative:* Minimize patients receiving inferior treatments
- *Exploration-exploitation* is life-or-death

- **A/B Testing**

- Test website designs, features, UI changes
- Each variant has unknown conversion rate
- Want to identify best variant quickly
- Minimize opportunity cost during testing

Applications: Real-World Decision-Making

- **Recommendation Systems**

- Netflix/Spotify: Which content to recommend?
- Each item has unknown appeal to user
- Must balance popular items vs. discovery
- Learn user preferences over time

- **Network Routing**

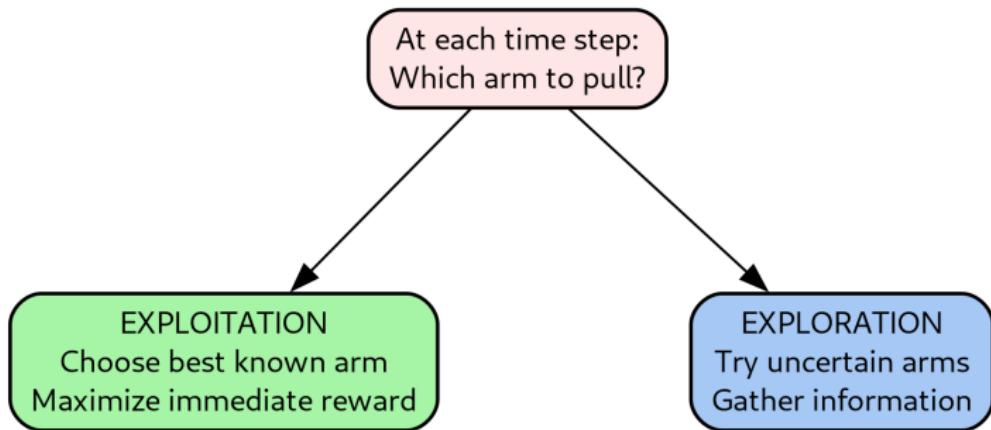
- Choose which server/path to route request
- Each option has variable latency/reliability
- Conditions change over time
- Need fast, adaptive decisions

- **Financial Portfolio Allocation**

- Allocate capital among investment options
- Each asset has unknown future returns
- Balance risk and reward
- Market conditions non-stationary

The Exploration–Exploitation Tradeoff

- The “Exploration–Exploitation Tradeoff” is the central dilemma of sequential learning
- **Why This Tradeoff is Unavoidable:**
 - *Pure exploitation*: Never try suboptimal-looking arms → may stick with suboptimal choice forever
 - *Pure exploration*: Always try random arms → collect lots of information but low total reward
 - *Must do both*: Explore enough to find good arms, exploit enough to gain high reward



The K-Armed Bandit Problem: Formal Definition

- **Setting:**

- K arms (actions), indexed $i = 1, \dots, K$
- Time horizon: T rounds (sometimes $T = \infty$)
- At round $t = 1, \dots, T$:
 1. Agent selects arm $A_t \in \{1, \dots, K\}$
 2. Environment reveals reward $R_t \sim P_{A_t}$
 3. Agent observes only R_t (not rewards from other arms)

- **Stochastic Bandit Assumption:**

- Each arm i has a fixed reward distribution P_i with mean $\mu_i = \mathbb{E}[R|A = i]$
- Rewards are independent across rounds: R_t iid
- The means μ_1, \dots, μ_K are unknown to the agent

- **Goal:** Design a policy (algorithm) that selects A_1, A_2, \dots, A_T to maximize

$$\text{Total Reward} = \sum_{t=1}^T R_t$$

- **Key Insight:** Can't maximize reward with certainty (rewards are random),
but can minimize *regret*



Reward Distributions

- Each arm i has an associated **reward distribution** P_i
 - When arm i is pulled, reward $R \sim P_i$ is sampled
 - Common distributions: Bernoulli, Gaussian, bounded support
- **Expected Reward** (Mean):
 - Each arm has true mean reward: $\mu_i = \mathbb{E}[R|A = i]$
 - This is the quantity we care about for maximizing total reward
 - The agent does not know μ_i initially
- **Unknown Parameters**: The Challenge
 - The means μ_1, \dots, μ_K are unknown
 - Also variance, shape of distribution may be unknown
 - Agent must estimate these from observed rewards
 - Must balance learning (exploration) and earning (exploitation)
- **Example**: Bernoulli Bandits
 - Each arm i returns reward $R_t = 1$ with probability μ_i , else $R_t = 0$
 - μ_i is the “success probability” or “conversion rate”
 - Common in A/B testing: μ_i = click-through rate of ad i

Performance Metrics

- How do we measure algorithm performance when rewards are random?
- **Optimal Arm:** Define $i^* = \operatorname{argmax}_i \mu_i$
 - The arm with highest expected reward
 - If we knew μ_1, \dots, μ_K , we would always pull arm i^*
 - Optimal expected reward per round: $\mu^* = \mu_{i^*}$
- **Instantaneous Regret** at time t :

$$\ell_t = \mu^* - \mu_{A_t}$$

- The difference between optimal mean and mean of chosen arm
 - How much expected reward we “lose” by not playing i^*
 - Note: $\ell_t \geq 0$ always
- **Cumulative Regret** over T rounds:

$$L_T = \sum_{t=1}^T \ell_t = T\mu^* - \sum_{t=1}^T \mu_{A_t}$$

- Total expected loss from suboptimal decisions
- Measures how much worse the algorithm is vs. always pulling i^*



Performance Metrics: Intuition

- **Why Regret?**

- Cannot maximize reward with certainty (randomness in rewards)
- But can minimize expected loss relative to optimal strategy
- Regret = “opportunity cost” of learning

- **Key Property:** Regret depends only on how many times we pull each arm

- Let $N_i(T) = \sum_{t=1}^T I[\{A_t = i\}]$ = number of times arm i pulled
- Then:

$$L_T = \sum_{i=1}^K (\mu^* - \mu_i) N_i(T) = \sum_{i=1}^K \Delta_i N_i(T)$$

- Where $\Delta_i = \mu^* - \mu_i$ is the “suboptimality gap” of arm i

- **Implications:**

- To minimize regret, pull suboptimal arms as few times as possible
- But we don't know which arms are suboptimal initially!
- Must explore to identify i^* , but exploration increases regret

Regret Lower Bounds

- **Why Lower Bounds Matter**

- Shows fundamental limits of *any* algorithm
- No algorithm can do better (up to constants)
- Tells us if an algorithm is near-optimal
- Guides algorithm design: what's achievable?

- **Intuition:** Why Can't Regret Be Constant?

- Need to distinguish between arms with close means
- If $\mu_1 \approx \mu_2$, need many samples to tell which is better
- Must pull suboptimal arms enough to be confident they're suboptimal
- Trade-off: confidence vs. regret

Regret Lower Bounds: Lai–Robbins Bound

- **Lai and Robbins (1985)** proved:
 - For any “consistent” algorithm (one that identifies optimal arm eventually)
 - The cumulative regret must satisfy:

$$\liminf_{T \rightarrow \infty} \frac{L_T}{\log T} \geq \sum_{i: \Delta_i > 0} \frac{\Delta_i}{D(\mu_i \| \mu^*)}$$

- Where $D(\mu_i \| \mu^*)$ is the KL divergence between reward distributions
- **Simplified Version** (for bounded rewards):
 - For any consistent algorithm:

$$L_T = \Omega \left(\sum_{i: \Delta_i > 0} \frac{\log T}{\Delta_i} \right)$$

- Regret must grow at least logarithmically with T
- **Interpretation:**
 - Cannot achieve constant regret (must be $\Omega(\log T)$)
 - Smaller gap $\Delta_i \rightarrow$ more pulls needed \rightarrow higher regret
 - Best achievable: $L_T = O(\log T)$ (logarithmic regret)

Regret Lower Bounds: Implications for Algorithm Design

- **Design Goal:** Achieve $O(\log T)$ regret
 - Matches lower bound (up to constants)
 - Called “order-optimal” or “asymptotically optimal”
- **Key Insight:** Must pull each suboptimal arm i roughly $O\left(\frac{\log T}{\Delta_i^2}\right)$ times
 - Arms with smaller gaps Δ_i need more exploration
 - Total regret: $\sum_i \Delta_i \cdot O\left(\frac{\log T}{\Delta_i^2}\right) = O\left(\frac{\log T}{\Delta_i}\right)$
- **What Won’t Work:**
 - Fixed exploration schedule (e.g., pull each arm n times): $\Omega(K)$ regret
 - Random sampling: $\Omega(\sqrt{T})$ regret
 - Need adaptive exploration that concentrates on promising arms

- Introduction
- ***Algorithms***
- Bayesian Approaches

Baseline Strategies

- Before studying sophisticated algorithms, consider naive approaches
- **Greedy Algorithm:**
 - At each time t :
 1. Compute empirical mean reward for each arm: $\hat{\mu}_i(t) = \frac{1}{N_i(t)} \sum_{\tau: A_\tau=i} R_\tau$
 2. Pull arm with highest empirical mean: $A_t = \operatorname{argmax}_i \hat{\mu}_i(t)$
 - Pure exploitation: always pull current best arm
- **Random Selection:**
 - At each time t : Choose A_t uniformly at random from $\{1, \dots, K\}$
 - Pure exploration: no use of observed rewards

Baseline Strategies: Limitations

- **Greedy Algorithm Fails:**

- *Problem:* May get “stuck” on suboptimal arm
- *Example:* Suppose arm 1 has $\mu_1 = 0.6$, arm 2 has $\mu_2 = 0.7$
 - If we happen to pull arm 1 first and get reward 1
 - Then $\hat{\mu}_1(1) = 1$, and we’ll keep pulling arm 1
 - Never try arm 2, which is actually better
- *Regret:* $L_T = \Omega(T)$ (linear regret, very bad!)

- **Random Selection Fails:**

- *Problem:* Never uses information from rewards
- Pulls all arms equally, even after learning which are bad
- *Regret:* $L_T = \Theta(\sqrt{KT})$ (grows with T , suboptimal)

- **Key Lesson:** Need to balance exploration and exploitation

- Greedy: too much exploitation
- Random: too much exploration
- Need adaptive strategy that explores less over time

ϵ -Greedy Algorithm

- **Simple Idea:** With small probability ϵ , explore; otherwise exploit
- **Algorithm:**
 - At each time t :
 1. Compute empirical means: $\hat{\mu}_i(t) = \frac{\sum_{\tau < t: A_\tau = i} R_\tau}{N_i(t)}$
 2. With probability ϵ : Choose arm uniformly at random (explore)
 3. With probability $1 - \epsilon$: Choose arm $\text{argmax}_i \hat{\mu}_i(t)$ (exploit)
 - **Hyperparameter:** $\epsilon \in (0, 1)$
 - Small ϵ (e.g., 0.01): mostly exploit, little exploration
 - Large ϵ (e.g., 0.5): explore half the time
 - Typical choice: $\epsilon = 0.1$

ϵ -Greedy: Exploration Schedules

- **Fixed ϵ :**
 - Use same ϵ for all t
 - *Problem:* Keeps exploring even after identifying best arm
 - *Regret:* $L_T = \Omega(T)$ (linear, not optimal)
- **Decaying ϵ_t :**
 - Use time-varying ϵ_t that decreases with t
 - Example: $\epsilon_t = \frac{1}{t}$ or $\epsilon_t = \min(1, \frac{K}{t})$
 - *Intuition:* Explore a lot early, less over time
 - *Regret:* Can achieve $L_T = O(T^{2/3})$ with right schedule (still not $O(\log T)$)
- **Limitations:**
 - Hard to choose optimal schedule without knowing Δ_i 's
 - Explores uniformly (doesn't focus on promising arms)
 - Better algorithms exist (UCB, Thompson Sampling)

ϵ -Greedy: Regret Behavior

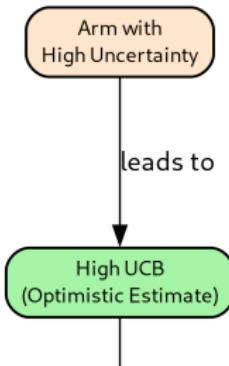
- **Fixed ϵ :** $L_T = \Theta(T)$
 - Explores ϵ fraction of time forever
 - Linear regret: $\epsilon T \sum_i \Delta_i / K$
- **Decaying $\epsilon_t = O(1/t^\alpha)$:**
 - If α too small: not enough exploration \rightarrow may miss i^*
 - If α too large: too much exploration \rightarrow high regret
 - Best tuned: $L_T = O(T^{2/3})$ (worse than $O(\log T)$ lower bound)
- **Practical Performance:**
 - Simple, easy to implement
 - Works reasonably well in practice
 - Often used as baseline
 - But theoretically suboptimal (cannot achieve $O(\log T)$)

Optimism in the Face of Uncertainty

- **Core Idea:** “Be optimistic about uncertain options”
 - If unsure about arm’s value, assume it could be the best
 - Optimism drives exploration of uncertain arms
 - As arms are explored, optimism diminishes (uncertainty shrinks)
- **Why Optimism Works:**
 - If an arm *is* optimal: Optimism ensures we keep trying it
 - If an arm *is not* optimal: We’ll eventually learn and stop pulling it
 - Natural exploration-exploitation balance without explicit randomization
- **Contrast with ϵ -Greedy:**
 - ϵ -greedy: Random exploration (ignores uncertainty)
 - Optimism: Directed exploration toward uncertain arms

Optimism: Confidence Intervals

- **Confidence Bound** for arm i after n pulls:
 - Empirical mean: $\hat{\mu}_i = \frac{1}{n} \sum_{\tau: A_\tau=i} R_\tau$
 - Uncertainty radius: $c_i(n) = \sqrt{\frac{2 \log T}{n}}$ (example)
 - Upper confidence bound: $UCB_i = \hat{\mu}_i + c_i(n)$
- **Interpretation:**
 - With high probability, true mean $\mu_i \leq UCB_i$
 - Fewer samples $n \rightarrow$ larger $c_i(n) \rightarrow$ more optimistic
 - More samples $n \rightarrow$ smaller $c_i(n) \rightarrow$ less optimistic
- **Optimistic Strategy:**
 - Pull arm with highest upper confidence bound: $A_t = \operatorname{argmax}_i UCB_i(t)$
 - Automatically explores uncertain arms and exploits promising arms



Upper Confidence Bound (UCB)

- **UCB1 Algorithm** (Auer, Cesa-Bianchi, Fischer, 2002):

- For first K rounds: Pull each arm once
- For round $t > K$:
 1. For each arm i , compute UCB:

$$\text{UCB}_i(t) = \hat{\mu}_i(t) + \sqrt{\frac{2 \log t}{N_i(t)}}$$

2. Pull arm: $A_t = \operatorname{argmax}_i \text{UCB}_i(t)$

- **Components:**

- $\hat{\mu}_i(t)$: Empirical mean (exploitation term)
- $\sqrt{\frac{2 \log t}{N_i(t)}}$: Exploration bonus (uncertainty term)
- $N_i(t)$: Number of times arm i pulled so far

Upper Confidence Bound: Arm Selection Rule

- **UCB Index:**

$$\text{UCB}_i(t) = \underbrace{\hat{\mu}_i(t)}_{\text{exploitation}} + \underbrace{\sqrt{\frac{2 \log t}{N_i(t)}}}_{\text{exploration bonus}}$$

- **How It Works:**

- Arms with high empirical mean $\hat{\mu}_i$ have high UCB (exploitation)
- Arms pulled few times (N_i small) have high UCB (exploration)
- Bonus $\propto \frac{1}{\sqrt{N_i}}$: Shrinks with pulls
- Bonus $\propto \sqrt{\log t}$: Grows slowly with time

- **Behavior Over Time:**

- *Early*: All arms have few pulls \rightarrow large bonuses \rightarrow explores
- *Middle*: Focuses on promising arms, occasionally revisits uncertain arms
- *Late*: Mostly pulls optimal arm (bonus becomes negligible)

Upper Confidence Bound: Key Intuition

- **Why UCB Works:**

- The bonus $\sqrt{\frac{2 \log t}{N_i(t)}}$ is carefully calibrated:
 - Derived from Hoeffding's concentration inequality
 - Ensures with high probability: $\mu_i \leq \text{UCB}_i(t)$
 - "Optimism": We act as if each arm is as good as its UCB

- **Self-Correcting:**

- If we pull a suboptimal arm too much: Its UCB decreases (more data \rightarrow less uncertainty)
- If we neglect the optimal arm: Its UCB increases ($\log t$ grows, while N_i stays fixed)
- Automatically balances exploration and exploitation

- **No Tuning Parameters:**

- Unlike ϵ -greedy (requires choosing ϵ)
- UCB has no hyperparameters (besides confidence constant)
- Adapts automatically to problem structure

Regret Analysis of UCB

- **Theorem** (Auer et al., 2002):
 - For UCB1 with rewards in $[0, 1]$, the cumulative regret satisfies:

$$L_T \leq 8 \sum_{i:\Delta_i > 0} \frac{\log T}{\Delta_i} + \left(1 + \frac{\pi^2}{3}\right) \sum_{i=1}^K \Delta_i$$

- **Simplified**: $L_T = O\left(\sum_{i:\Delta_i > 0} \frac{\log T}{\Delta_i}\right)$
- **Interpretation**:
 - Logarithmic regret: $L_T = O(\log T)$
 - Matches Lai–Robbins lower bound (up to constants)
 - *Order-optimal*: Cannot do better asymptotically
- **Dependence on Δ_i** :
 - Smaller gaps $\Delta_i \rightarrow$ higher regret (harder to distinguish arms)
 - Arms with $\Delta_i = 0$ (i.e., optimal arms) don't contribute to regret bound

Regret Analysis: Role of Concentration Inequalities

- **Key Tool:** Hoeffding's Inequality
 - For rewards R_1, \dots, R_n iid in $[0, 1]$ with mean μ :

$$\Pr(|\hat{\mu} - \mu| \geq \epsilon) \leq 2 \exp(-2n\epsilon^2)$$

- **Application to UCB:**

- Choose radius $c = \sqrt{\frac{2 \log t}{n}}$ so that:

$$\Pr\left(\mu_i > \hat{\mu}_i + \sqrt{\frac{2 \log t}{N_i(t)}}\right) \leq \frac{1}{t^4}$$

- With high probability, the true mean is below the UCB
- **Union Bound:** Across all arms and time steps:
 - Probability that any UCB is violated at any time $\leq \sum_{t=1}^T \sum_{i=1}^K \frac{1}{t^4} < \infty$
 - UCB indices are “optimistic” with high probability
- **Consequence:** If optimal arm has high UCB, we'll pull it; if suboptimal arm has high UCB, we'll pull it until UCB decreases

Regret Comparison: UCB vs. ϵ -Greedy

ϵ -Greedy

- **Regret:** $L_T = O(T)$ (fixed ϵ) or $O(T^{2/3})$ (decaying ϵ)
- **Exploration:** Uniform random sampling
- **Pros:**
 - Simple to implement
 - Easy to understand
- **Cons:**
 - Linear or polynomial regret
 - Ignores uncertainty
 - Requires tuning ϵ
- **Practical Takeaway:** UCB is theoretically and empirically superior for stationary problems

UCB

- **Regret:** $L_T = O(\log T)$
- **Exploration:** Directed by uncertainty
- **Pros:**
 - Logarithmic regret (optimal)
 - No hyperparameters
 - Adaptive
- **Cons:**
 - Slightly more complex
 - Requires confidence bounds

Variants of UCB

- **UCB1**: Original algorithm (assumes bounded rewards in $[0, 1]$)
- **UCB-V** (Variance-Aware UCB):

- Incorporates empirical variance of rewards
- Index:

$$\text{UCB-V}_i(t) = \hat{\mu}_i(t) + \sqrt{\frac{2\hat{\sigma}_i^2(t) \log t}{N_i(t)}} + \frac{3 \log t}{N_i(t)}$$

- Where $\hat{\sigma}_i^2(t)$ is sample variance of arm i
- **Benefit**: Better for arms with low variance (tighter bounds)

- **KL-UCB**:

- Uses Kullback–Leibler divergence for tighter confidence bounds
- Index:

$$\text{KL-UCB}_i(t) = \max \{q : N_i(t) D(\hat{\mu}_i(t) \| q) \leq \log t\}$$

- Where $D(p \| q)$ is KL divergence
- **Benefit**: Asymptotically optimal constants (better than UCB1)

Variants of UCB: When They Matter

- **UCB1:** Good default choice
 - Simple, no tuning, strong guarantees
 - Works for any reward distribution with bounded support
- **UCB-V:** Use when variance varies across arms
 - Example: Some ads have consistent CTR, others highly variable
 - Can reduce regret if some arms have low variance
- **KL-UCB:** Use when asymptotic performance matters
 - Better constants in regret bound
 - More computationally expensive (requires solving equation for index)
 - Recommended for Bernoulli or Gaussian bandits
- **Practical Advice:** Start with UCB1; switch to variants if:
 - You have domain knowledge about variance structure (UCB-V)
 - You need best possible asymptotic performance (KL-UCB)

- Introduction
- Algorithms
- *Bayesian Approaches*

Bayesian Bandits

- So far: **Frequentist** approach (UCB, ϵ -greedy)
 - Treat μ_i as fixed unknown constants
 - Use confidence bounds and empirical means
- **Bayesian Approach:** Treat μ_i as random variables
 - Start with prior distribution $p(\mu_i)$ for each arm
 - Update beliefs using Bayes' rule after each observation
 - Make decisions based on posterior distributions
- **Bayesian Framework:**
 - Prior: $p(\mu_i)$ (e.g., $\mu_i \sim \text{Beta}(1, 1)$ for Bernoulli arm)
 - Likelihood: $p(R|\mu_i)$ (e.g., $R \sim \text{Bernoulli}(\mu_i)$)
 - Posterior: $p(\mu_i|R_{1:t}) \propto p(R_{1:t}|\mu_i)p(\mu_i)$ (by Bayes' rule)

Bayesian Bandits: Priors and Posteriors

- **Prior Distribution:** Encodes initial beliefs about μ_i
 - Uniform prior: $\mu_i \sim \text{Beta}(1, 1)$ (no prior knowledge)
 - Informative prior: $\mu_i \sim \text{Beta}(\alpha, \beta)$ (domain knowledge)
- **Posterior Update:** After observing rewards from arm i
 - Suppose arm i has been pulled n times: s successes, f failures
 - Posterior (for Bernoulli arm with Beta prior):
$$p(\mu_i | \text{data}) = \text{Beta}(\alpha + s, \beta + f)$$
- Conjugate prior: Posterior has same form as prior (easy to update)
- **Bayesian Inference:** Posterior encodes all information about μ_i
 - Mean: $\mathbb{E}[\mu_i | \text{data}] = \frac{\alpha+s}{\alpha+\beta+n}$
 - Uncertainty: $\mathbb{V}[\mu_i | \text{data}] = \frac{(\alpha+s)(\beta+f)}{(\alpha+\beta+n)^2(\alpha+\beta+n+1)}$
 - More data \rightarrow sharper posterior (less uncertainty)

Bayesian Decision-Making

- **How to Choose Arm?** Several strategies:
- **Greedy Bayesian:** Choose arm with highest posterior mean

$$A_t = \operatorname{argmax}_i \mathbb{E}[\mu_i | \text{data}_t]$$

- Problem: No exploration (same issue as frequentist greedy)
 - **Bayes-UCB:** Use posterior quantile as UCB
- $$A_t = \operatorname{argmax}_i Q_i\left(1 - \frac{1}{t}\right)$$
- Where $Q_i(p)$ is the p -th quantile of posterior $p(\mu_i | \text{data}_t)$
 - Similar to UCB, but uses Bayesian uncertainty
- **Probability of Best Arm:** Choose arm most likely to be optimal

$$A_t = \operatorname{argmax}_i \Pr(\mu_i = \max_j \mu_j | \text{data}_t)$$

- **Thompson Sampling:** Sample from posteriors and play best sample

(next slide)

Thompson Sampling

- **Posterior Sampling** (Thompson, 1933; Agrawal & Goyal, 2012)
 - At each time t :
 1. Sample $\tilde{\mu}_i \sim p(\mu_i|\text{data}_t)$ for each arm i
 2. Pull arm $A_t = \operatorname{argmax}_i \tilde{\mu}_i$
 3. Observe reward R_t
 4. Update posterior for arm A_t
- **Intuition:** “Probability matching”
 - Pull arm i with probability $\approx \Pr(i \text{ is optimal}|\text{data})$
 - If arm likely to be best, sample will often be highest
 - If arm uncertain, might still be sampled (exploration)
- **Example:** Bernoulli Bandits with Beta Priors
 - Posterior: $p(\mu_i|\text{data}) = \text{Beta}(\alpha_i, \beta_i)$
 - Algorithm:
 1. Sample $\tilde{\mu}_i \sim \text{Beta}(\alpha_i, \beta_i)$ for each i
 2. Pull $A_t = \operatorname{argmax}_i \tilde{\mu}_i$
 3. If reward $R_t = 1$: $\alpha_{A_t} \leftarrow \alpha_{A_t} + 1$; else $\beta_{A_t} \leftarrow \beta_{A_t} + 1$

Thompson Sampling: Algorithm Steps

Thompson Sampling for Bernoulli Bandits

Input: Number of arms K , time horizon T

Initialize: For each arm $i = 1, \dots, K$: - $\alpha_i \leftarrow 1, \beta_i \leftarrow 1$ (Beta(1,1) prior)

For $t = 1, 2, \dots, T$:

1. **Sample** from posteriors:
 - For each arm i : $\tilde{\mu}_i \sim \text{Beta}(\alpha_i, \beta_i)$
2. **Select arm**: $A_t \leftarrow \text{argmax}_i \tilde{\mu}_i$
3. **Observe reward**: R_t
4. **Update posterior** for arm A_t :
 - If $R_t = 1$: $\alpha_{A_t} \leftarrow \alpha_{A_t} + 1$
 - If $R_t = 0$: $\beta_{A_t} \leftarrow \beta_{A_t} + 1$

Thompson Sampling: Practical Performance

- **Theoretical Guarantees:**

- Achieves $L_T = O(\log T)$ regret (optimal)
- Matches UCB asymptotically
- Proofs more involved than UCB (Agrawal & Goyal, 2012)

- **Empirical Performance:**

- Often *outperforms* UCB in practice (better constants)
- Faster convergence to optimal arm in many problems
- More robust to model misspecification

- **Advantages:**

- Naturally incorporates prior knowledge (via priors)
- Flexible: Works with any reward distribution (if conjugate prior available)
- Simple randomized algorithm (no complex optimization)

- **Disadvantages:**

- Requires choosing prior (though often uniform prior works well)
- Posterior updates may be non-trivial for complex distributions
- Theoretically less understood than UCB (until recent years)

- **Current Status:** *De facto* standard in practice (especially industry)

Comparison of UCB and Thompson Sampling

- Frequentist vs. Bayesian views
- Empirical vs. theoretical tradeoffs
- When to use each

Adversarial Bandits

- Motivation for adversarial models
- Difference from stochastic bandits
- Adversary assumptions

EXP3 Algorithm

- Probability-weighted arm selection
- Importance weighting
- Regret guarantees

Stochastic vs. Adversarial Settings

- Model assumptions
- Algorithm robustness
- Use cases

Contextual Bandits

- Using side information
- Feature vectors
- Decision rules with context

Linear Contextual Bandits

- Linear reward model
- Parameter estimation
- Assumptions

- Algorithm structure
- Confidence ellipsoids
- Regret guarantees

Contextual Thompson Sampling

- Bayesian linear models
- Posterior sampling with context
- Comparison to LinUCB

Bandits vs. Reinforcement Learning

- Stateless vs. stateful problems
- Bandits as a special case of RL
- When RL is needed

Non-Stationary Bandits

- Changing reward distributions
- Concept drift
- Practical motivation

Algorithms for Non-Stationary Bandits

- Discounted UCB
- Sliding-window methods
- Restart strategies

Structured Bandits

- Linear bandits
- Combinatorial bandits
- Graph and rank-one structures

Best-Arm Identification

- Pure exploration setting
- Fixed-confidence vs. fixed-budget
- Sample complexity

Practical Issues

- Delayed feedback
- Offline evaluation
- Safety constraints